

# week2\_premium\_estimator\_main

June 23, 2025

```
[5]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split, KFold, \
    cross_val_score, cross_validate, RandomizedSearchCV
from sklearn.metrics import mean_squared_error, r2_score, root_mean_squared_error
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from time import time
```

## 0.1 Data Import & Exploration

```
[6]: # Switching to the source directory

parent_dir = os.path.abspath(os.path.join(os.getcwd(), "../.."))
parent_dir
print(parent_dir)
```

C:\Users\91948\Downloads\BKs\Projects\personalized\_health\_insurance\_premium\_estimator

```
[7]: # Reading the data

df = pd.read_excel(fr'{parent_dir}\data\premiums.xlsx', sheet_name='Sheet1')
```

```
[8]: # Displaying the data

df.head()
```

```
[8]:   Age  Gender  Region  Marital_status  Number Of Dependants  BMI_Category  \
0   26   Male  Northwest      Unmarried                      0      Normal
1   29  Female  Southeast      Married                      2      Obesity
2   49  Female  Northeast      Married                      2      Normal
```

3	30	Female	Southeast	Married	3	Normal
4	18	Male	Northeast	Unmarried	0	Overweight

	Smoking_Status	Employment_Status	Income_Level	Income_Lakhs	\
0	No Smoking	Salaried	<10L	6	
1	Regular	Salaried	<10L	6	
2	No Smoking	Self-Employed	10L - 25L	20	
3	No Smoking	Salaried	> 40L	77	
4	Regular	Self-Employed	> 40L	99	

	Medical History	Insurance_Plan	Annual_Premium_Amount
0	Diabetes	Bronze	9053
1	Diabetes	Bronze	16339
2	High blood pressure	Silver	18164
3	No Disease	Gold	20303
4	High blood pressure	Silver	13365

```
[9]: df.shape
```

```
[9]: (50000, 13)
```

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   50000 non-null  int64
1   Gender                               50000 non-null  object
2   Region                               50000 non-null  object
3   Marital_status                       50000 non-null  object
4   Number Of Dependants                 50000 non-null  int64
5   BMI_Category                         50000 non-null  object
6   Smoking_Status                       49989 non-null  object
7   Employment_Status                   49998 non-null  object
8   Income_Level                         49987 non-null  object
9   Income_Lakhs                         50000 non-null  int64
10  Medical History                      50000 non-null  object
11  Insurance_Plan                       50000 non-null  object
12  Annual_Premium_Amount                50000 non-null  int64
dtypes: int64(4), object(9)
memory usage: 5.0+ MB
```

## 0.2 Data Cleaning

The following steps were performed during data cleaning:

- Handling missing values

- Removing duplicated rows
- Formatting numeric and categorical values
- Treating outliers to improve data quality

### 0.2.1 Column Formatter

```
[11]: # Renaming the columns with proper formatter
# Eg: Number Of Dependants -> number_of_dependants

df.columns = df.columns.str.replace(' ', '_').str.lower()
```

```
[12]: df.head()
```

```
[12]:   age  gender  region marital_status  number_of_dependants  bmi_category \
0   26   Male  Northwest      Unmarried                0      Normal
1   29  Female  Southeast      Married                2      Obesity
2   49  Female  Northeast      Married                2      Normal
3   30  Female  Southeast      Married                3      Normal
4   18   Male  Northeast      Unmarried                0  Overweight
```

```
   smoking_status  employment_status  income_level  income_lakhs \
0      No Smoking           Salaried      <10L             6
1           Regular           Salaried      <10L             6
2      No Smoking  Self-Employed  10L - 25L            20
3      No Smoking           Salaried      > 40L            77
4           Regular  Self-Employed      > 40L            99
```

```
   medical_history  insurance_plan  annual_premium_amount
0           Diabetes           Bronze             9053
1           Diabetes           Bronze            16339
2  High blood pressure           Silver            18164
3           No Disease            Gold            20303
4  High blood pressure           Silver            13365
```

### 0.2.2 Handling Missing Values

```
[13]: # Finding the columns with NA values

df.isna().sum()
```

```
[13]: age                0
gender                0
region                0
marital_status        0
number_of_dependants  0
bmi_category          0
smoking_status        11
employment_status      2
```

```

income_level      13
income_lakhs      0
medical_history    0
insurance_plan     0
annual_premium_amount  0
dtype: int64

```

```
[14]: # Dropping the Rows where it has NA value
```

```

# Before Dropping
df[df['smoking_status'].isna()].head()

```

```

[14]:      age  gender      region marital_status  number_of_dependants \
177     26  Female  Southwest      Married              3
15648   47   Male  Southwest      Married              4
16324   45   Male  Northwest      Married              4
16941   34   Male  Southwest      Married              5
16975   23   Male  Southwest      Unmarried             0

      bmi_category smoking_status employment_status income_level \
177    Underweight          NaN      Salaried      > 40L
15648      Normal          NaN      Freelancer  25L - 40L
16324   Overweight          NaN      Salaried  10L - 25L
16941      Normal          NaN  Self-Employed  25L - 40L
16975      Normal          NaN      Freelancer    <10L

      income_lakhs      medical_history insurance_plan \
177             69             Diabetes          Gold
15648            32             Diabetes          Gold
16324            16  High blood pressure & Heart disease  Silver
16941            35  High blood pressure & Heart disease    Gold
16975             3             No Disease        Bronze

      annual_premium_amount
177                22605
15648               26100
16324               21881
16941               25865
16975                6001

```

Since the count is low, the rows containing at least one missing value were dropped.

```

[15]: # Dropping
df.dropna(how='any',inplace = True)

```

```

[16]: # After Dropping
df[df['smoking_status'].isna()]

```

```
[16]: Empty DataFrame
      Columns: [age, gender, region, marital_status, number_of_dependants,
      bmi_category, smoking_status, employment_status, income_level, income_lakhs,
      medical_history, insurance_plan, annual_premium_amount]
      Index: []
```

```
[17]: df.isna().sum()
```

```
[17]: age                0
      gender            0
      region            0
      marital_status    0
      number_of_dependants  0
      bmi_category      0
      smoking_status    0
      employment_status  0
      income_level      0
      income_lakhs      0
      medical_history   0
      insurance_plan    0
      annual_premium_amount  0
      dtype: int64
```

```
[18]: df.reset_index(inplace=True,drop=True)
```

### 0.2.3 Handling Duplicated Rows

```
[19]: # Duplicated rows

      df[df.duplicated()]
```

```
[19]: Empty DataFrame
      Columns: [age, gender, region, marital_status, number_of_dependants,
      bmi_category, smoking_status, employment_status, income_level, income_lakhs,
      medical_history, insurance_plan, annual_premium_amount]
      Index: []
```

Although there are no duplicated rows, they are being dropped as a precautionary measure.

```
[20]: # Dropping the duplicated rows

      df.drop_duplicates(inplace=True)
```

### 0.2.4 Formatting Values - Numeric Columns

```
[21]: df.dtypes
```

```
[21]: age                int64
      gender            object
      region            object
      marital_status    object
      number_of_dependants  int64
      bmi_category      object
      smoking_status    object
      employment_status object
      income_level      object
      income_lakhs      int64
      medical_history   object
      insurance_plan    object
      annual_premium_amount  int64
      dtype: object
```

```
[22]: # Selecting only the columns with numerical values
```

```
numeric_columns = df.select_dtypes(include=['int64']).columns
numeric_columns
```

```
[22]: Index(['age', 'number_of_dependants', 'income_lakhs', 'annual_premium_amount'],
      dtype='object')
```

```
[23]: # Printing the unique values of each numeric columns to identify the values
      ↪with improper formats
```

```
for col in numeric_columns:
    print(f'{col}:\n',df[col].unique())
    print('*'*100)
```

age:

```
[ 26  29  49  30  18  56  33  43  59  22  21  46  68  60  27  25  36  20
  28  32  19  55  35  52  40  23  50  41  67  37  24  34  54  42  45  44
  57  38  31  58  48  51 224  47  39  53  66  64  65  62  61  70  72  69
  71 124  63 136 203 356]
```

```
*****
*****
```

number\_of\_dependants:

```
[ 0  2  3  4  1  5 -3 -1]
```

```
*****
*****
```

income\_lakhs:

```
[  6  20  77  99  14   4  46  21   3  97   1  27  15  18   7  37  30  13
   8  83  19  29   5  70  11  33  23  40  84  22   9  71  59  38  35  28
  39  57  25  12  36  92   2  24  16  34  93  78  26  49  68  52  62  31
  90  50  32  10  88  54  86  41  95  64  85  81  79  56  80  17  98  89
  82 100  44  66  53  75  94  69  58  74  65  91  42  61  87  96  51  67
  43  73  63  55  48  45  47  72  60 560  76 440 630 900 930 580 700 790]
```

```

770 680]
*****
*****
annual_premium_amount:
[ 9053 16339 18164 ... 26370 10957 27076]
*****
*****

```

The negative values in `number_of_dependents` should be handled by converting them to their absolute values.

[24]: *# Before formatting*

```
df['number_of_dependants'].unique()
```

[24]: array([ 0, 2, 3, 4, 1, 5, -3, -1])

[25]: `df['number_of_dependants'] = abs(df['number_of_dependants'])`

[26]: *# After formatting*

```
df['number_of_dependants'].unique()
```

[26]: array([0, 2, 3, 4, 1, 5])

## 0.2.5 Formatting Values - Categorical Columns

[27]: *# Selecting only the columns with categorical values*

```
cat_columns = df.select_dtypes(include=['object']).columns
cat_columns
```

[27]: Index(['gender', 'region', 'marital\_status', 'bmi\_category', 'smoking\_status',  
'employment\_status', 'income\_level', 'medical\_history',  
'insurance\_plan'],  
dtype='object')

[28]: *# Printing the unique values of each categorical columns to identify the values, ↵  
↵with improper formats*

```
for col in cat_columns:
    print(f'{col}:\n',list(df[col].unique()))
    print('*'*100)
```

gender:

```
['Male', 'Female']
```

```
*****
*****
```

region:

```

['Northwest', 'Southeast', 'Northeast', 'Southwest']
*****
*****
marital_status:
['Unmarried', 'Married']
*****
*****
bmi_category:
['Normal', 'Obesity', 'Overweight', 'Underweight']
*****
*****
smoking_status:
['No Smoking', 'Regular', 'Occasional', 'Smoking=0', 'Does Not Smoke', 'Not
Smoking']
*****
*****
employment_status:
['Salaried', 'Self-Employed', 'Freelancer']
*****
*****
income_level:
['<10L', '10L - 25L', '> 40L', '25L - 40L']
*****
*****
medical_history:
['Diabetes', 'High blood pressure', 'No Disease', 'Diabetes & High blood
pressure', 'Thyroid', 'Heart disease', 'High blood pressure & Heart disease',
'Diabetes & Thyroid', 'Diabetes & Heart disease']
*****
*****
insurance_plan:
['Bronze', 'Silver', 'Gold']
*****
*****

```

As indicated by the above result, the `smoking_status` column contains multiple values with the same meaning; therefore, this column will be formatted.

```
[29]: # Before formatting
```

```
df['smoking_status'].unique()
```

```
[29]: array(['No Smoking', 'Regular', 'Occasional', 'Smoking=0',
          'Does Not Smoke', 'Not Smoking'], dtype=object)
```

```
[30]: # Replacing values with desired formats
```

```
df['smoking_status'] = df['smoking_status'].replace(
```



```
{
    'Smoking=0' : 'No Smoking',
    'Does Not Smoke' : 'No Smoking',
    'Not Smoking' : 'No Smoking',
}
```

```
[31]: # After formatting
```

```
df['smoking_status'].unique()
```

```
[31]: array(['No Smoking', 'Regular', 'Occasional'], dtype=object)
```

### 0.2.6 Outlier Treatment

```
[32]: df.describe()
```

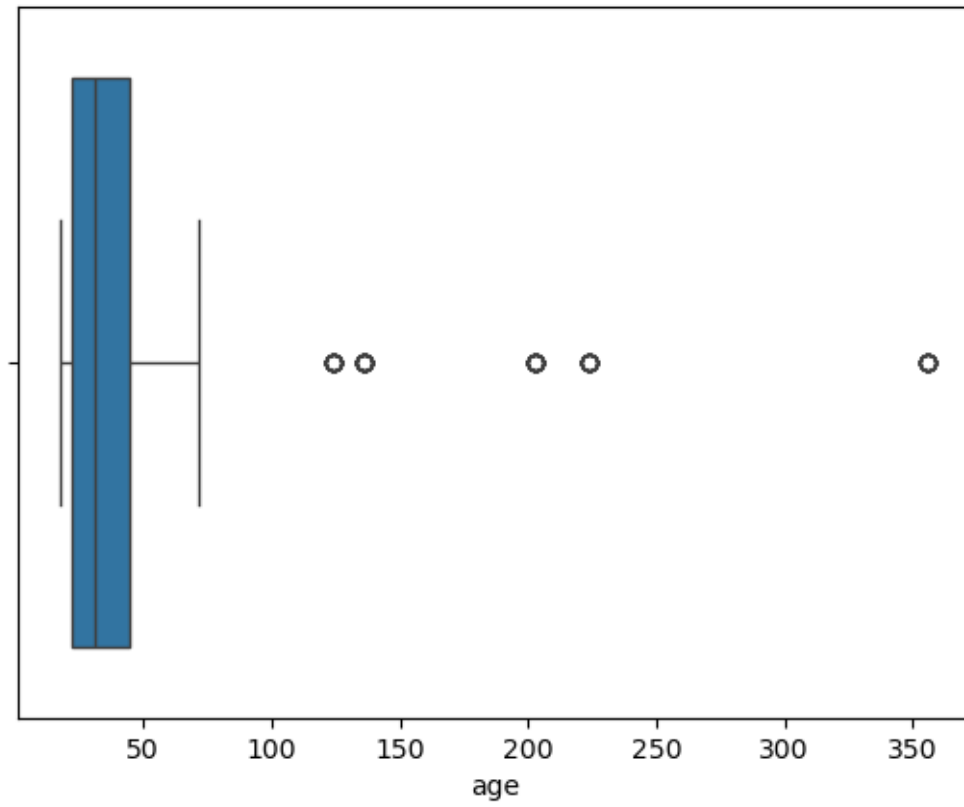
```
[32]:
```

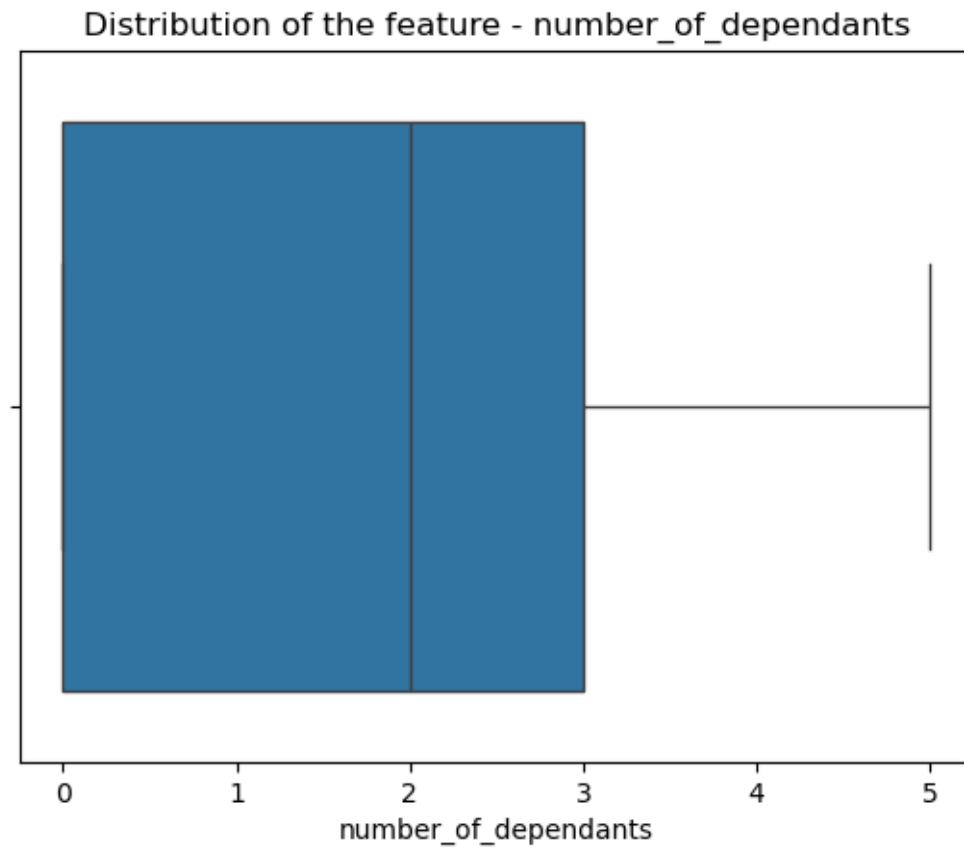
	age	number_of_dependants	income_lakhs	annual_premium_amount
count	49976.000000	49976.000000	49976.000000	49976.000000
mean	34.591764	1.717284	23.021150	15766.810189
std	15.000378	1.491953	24.221794	8419.995271
min	18.000000	0.000000	1.000000	3501.000000
25%	22.000000	0.000000	7.000000	8607.750000
50%	31.000000	2.000000	17.000000	13928.000000
75%	45.000000	3.000000	31.000000	22273.500000
max	356.000000	5.000000	930.000000	43471.000000

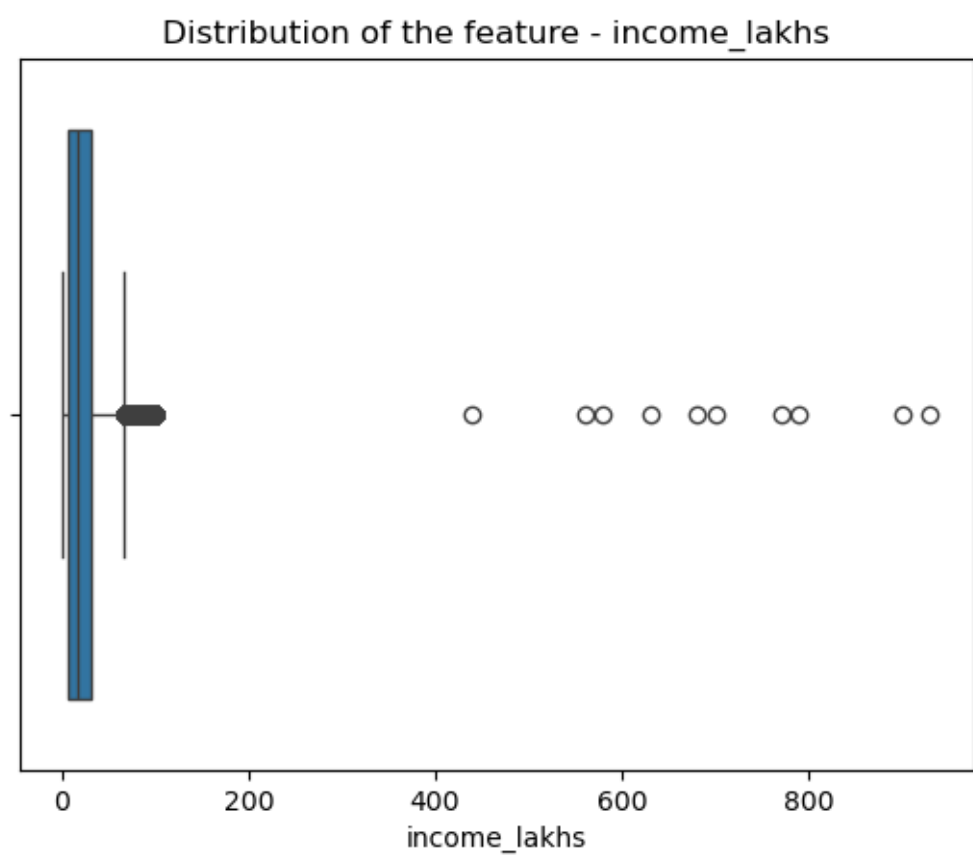
```
[33]: # Plotting box plot for each numerical columns to detect outliers
```


```
for col in numeric_columns:
    sns.boxplot(data=df,x=col)
    plt.title(f'Distribution of the feature - {col}')
    plt.show()
```

Distribution of the feature - age









As observed above, outliers are present in the `age` and `income_lakhs` columns. This is also illustrated by the box plot shown above.

### Age

```
[34]: # Before
      df['age'].unique()
```

```
[34]: array([ 26,  29,  49,  30,  18,  56,  33,  43,  59,  22,  21,  46,  68,
          60,  27,  25,  36,  20,  28,  32,  19,  55,  35,  52,  40,  23,
          50,  41,  67,  37,  24,  34,  54,  42,  45,  44,  57,  38,  31,
          58,  48,  51, 224,  47,  39,  53,  66,  64,  65,  62,  61,  70,
          72,  69,  71, 124,  63, 136, 203, 356])
```

```
[35]: # Selecting only the rows where age <= 100. Because age > 100 is outlier

      df1 = df[df['age'] <= 100]
      df1.reset_index(inplace=True,drop=True)
```

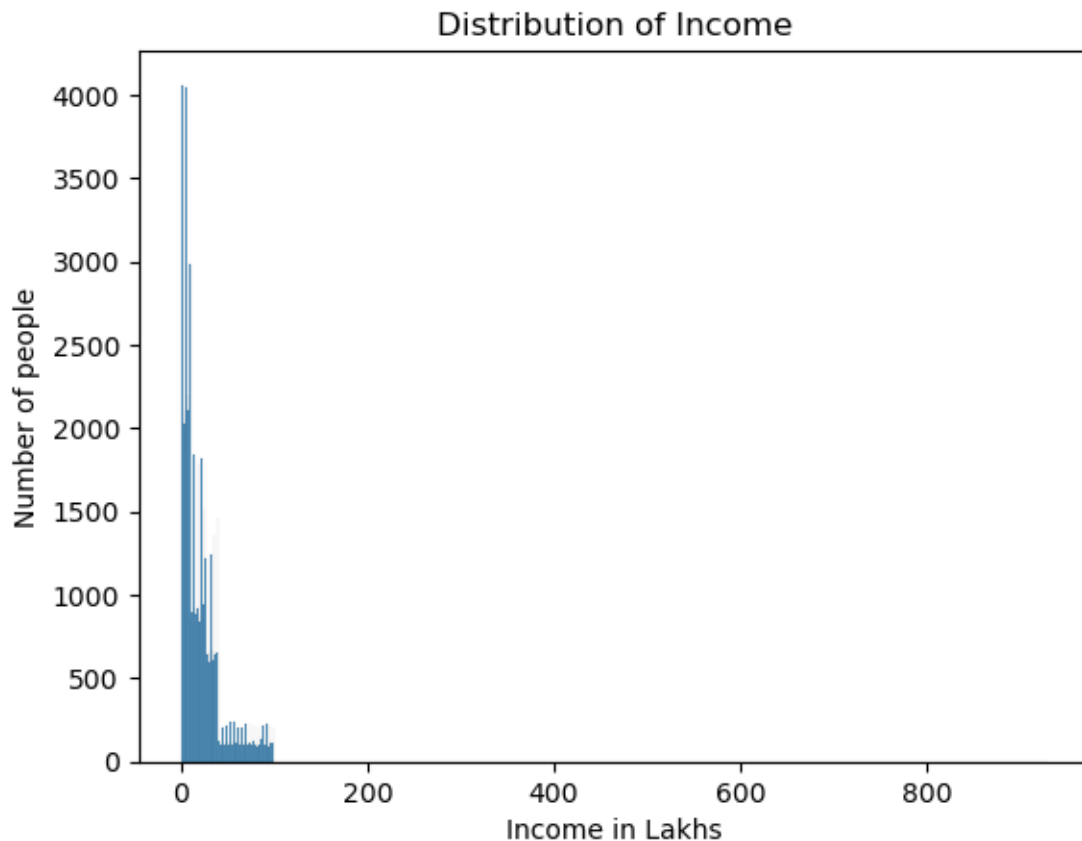
```
[36]: # After
      df1['age'].unique()
```

```
[36]: array([26, 29, 49, 30, 18, 56, 33, 43, 59, 22, 21, 46, 68, 60, 27, 25, 36,
          20, 28, 32, 19, 55, 35, 52, 40, 23, 50, 41, 67, 37, 24, 34, 54, 42,
          45, 44, 57, 38, 31, 58, 48, 51, 47, 39, 53, 66, 64, 65, 62, 61, 70,
          72, 69, 71, 63])
```

### Income

```
[37]: # Distribution of Income using Histogram
```

```
sns.histplot(data=df1,x='income_lakhs')
plt.title('Distribution of Income')
plt.xlabel('Income in Lakhs')
plt.ylabel('Number of people')
plt.show()
```



```
[38]: # Determining Lower Salary Boundary and Upper Salary Boundary using IQR Method
```

```
# Quartile 1 and Quartile 3
```

```
q1 = np.percentile(df1['income_lakhs'],25)
q3 = np.percentile(df1['income_lakhs'],75)
```

```

print(f'Q1 -> {q1}\nQ3 -> {q3}')

# Inter Quartile Range

iqr = q3-q1
print('IQR ->',iqr)

# Lower and Upper boundary using IQR

lower_boundary = q1 - (iqr * 1.5)
upper_boundary = q3 + (iqr * 1.5)
print(f'Lower Boundary -> {lower_boundary}\nUpper Boundary -> {upper_boundary}')

```

```

Q1 -> 7.0
Q3 -> 31.0
IQR -> 24.0
Lower Boundary -> -29.0
Upper Boundary -> 67.0

```

NOTE:

The current upper boundary is too low. Therefore, consultation with the business team will be conducted to determine the optimal upper boundary for identifying outliers in the income\_lakhs column.

It has been decided that the upper boundary will be set at 1 crore (i.e., 100 lakhs). Any values exceeding this threshold will be considered outliers.

```

[39]: # Selecting only the rows where income <= 100 lakhs. Because income > 100 lakhs
      ↳ is outlier

```

```

income_threshold = 100
df2 = df1[df1['income_lakhs'] <= income_threshold]

```

```

[40]: df2.reset_index(drop=True,inplace=True)

```

```

[41]: df2

```

```

[41]:
   age  gender  region marital_status  number_of_dependants  \
0    26   Male  Northwest      Unmarried                    0
1    29  Female  Southeast        Married                    2
2    49  Female  Northeast        Married                    2
3    30  Female  Southeast        Married                    3
4    18   Male  Northeast      Unmarried                    0
...   ...   ...   ...             ...                     ...
49903  24  Female  Northwest      Unmarried                    0
49904  47  Female  Southeast        Married                    2
49905  21   Male  Northwest      Unmarried                    0
49906  18   Male  Northwest      Unmarried                    2

```

49907	48	Female	Southwest	Married	3
-------	----	--------	-----------	---------	---

	bmi_category	smoking_status	employment_status	income_level \
0	Normal	No Smoking	Salaried	<10L
1	Obesity	Regular	Salaried	<10L
2	Normal	No Smoking	Self-Employed	10L - 25L
3	Normal	No Smoking	Salaried	> 40L
4	Overweight	Regular	Self-Employed	> 40L
...	...	...	...	...
49903	Underweight	No Smoking	Self-Employed	25L - 40L
49904	Normal	No Smoking	Salaried	> 40L
49905	Normal	Regular	Freelancer	25L - 40L
49906	Normal	No Smoking	Salaried	10L - 25L
49907	Normal	Occasional	Self-Employed	<10L

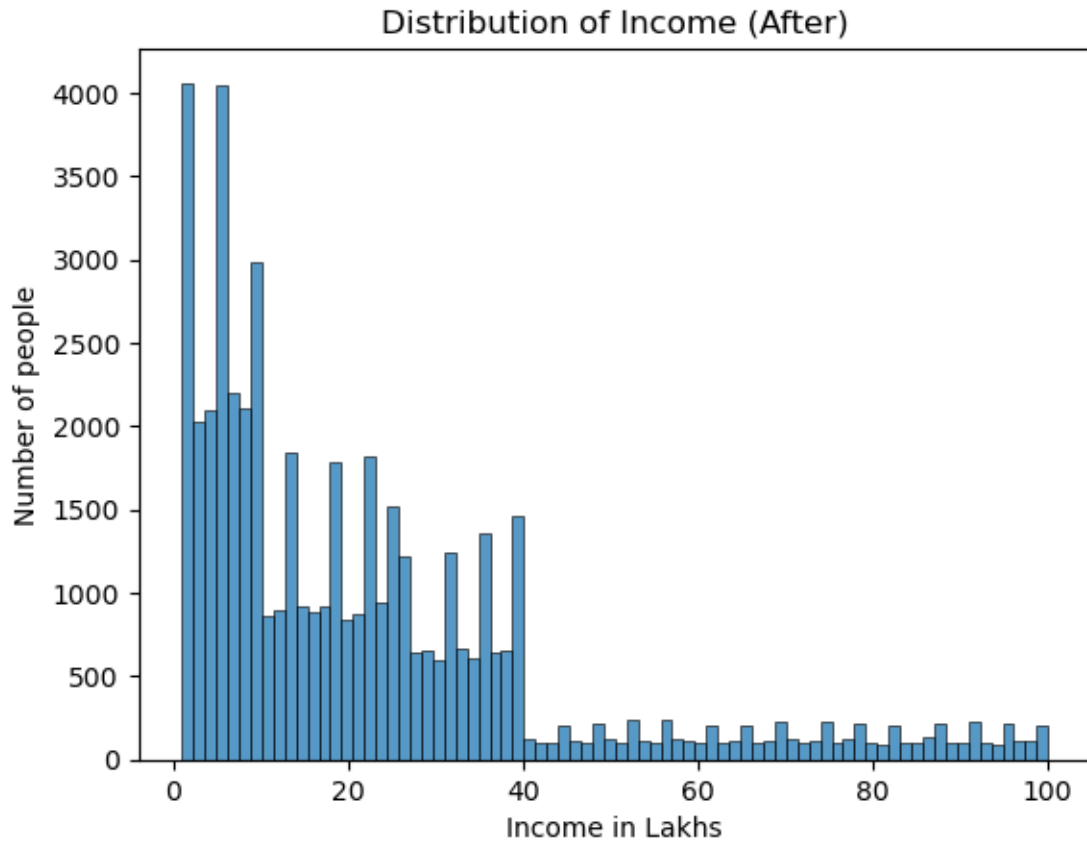
	income_lakhs	medical_history	insurance_plan	annual_premium_amount
0	6	Diabetes	Bronze	9053
1	6	Diabetes	Bronze	16339
2	20	High blood pressure	Silver	18164
3	77	No Disease	Gold	20303
4	99	High blood pressure	Silver	13365
...	...	...	...	...
49903	35	No Disease	Bronze	9111
49904	82	Thyroid	Gold	27076
49905	32	No Disease	Bronze	8564
49906	20	No Disease	Bronze	9490
49907	7	Diabetes	Silver	19730

```
[49908 rows x 13 columns]
```

```
[42]: # After Treating outliers in 'income_lakhs' columns
```

```
sns.histplot(data=df2,x='income_lakhs')
plt.title('Distribution of Income (After)')
plt.xlabel('Income in Lakhs')
plt.ylabel('Number of people')
plt.show()
```





### 0.3 EDA

Exploratory Data Analysis (EDA) is performed as follows:

- Univariate analysis on numeric columns
- Univariate analysis on categorical columns
- Bivariate analysis on numeric columns
- Bivariate analysis on categorical columns

#### 0.3.1 Univariate Analysis - Numeric Columns

The distribution of numerical columns will be plotted using histograms.

```
[43]: numeric_columns
```

```
[43]: Index(['age', 'number_of_dependants', 'income_lakhs', 'annual_premium_amount'],
dtype='object')
```

```
[44]: # Plotting Distribution for all numerical columns
```

```
fig , ax = plt.subplots(2,2,figsize=(8,8))
```

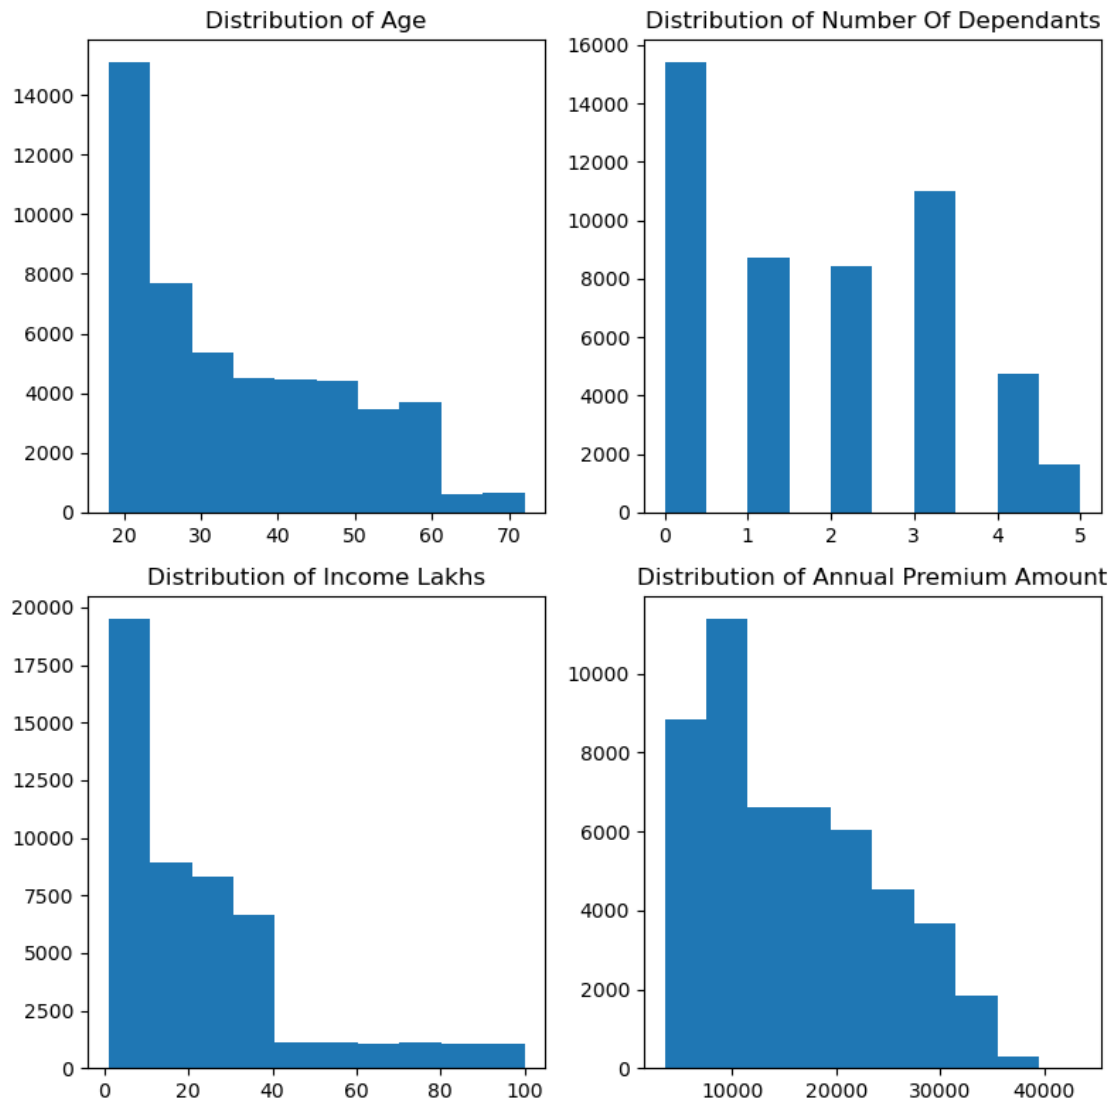
```

# Iterator initiation to retrieve one value at a time
it = iter(numeric_columns)

for i in range(2):
    for j in range(2):
        # To retrieve one value at a time
        col = next(it)
        # Format column name for display: capitalize first letter of each word
        ↪and replace underscores with spaces
        x_ax = col.title().replace('_', ' ')
        ax[i,j].hist(x=df2[col])
        ax[i,j].set_title(f'Distribution of {x_ax}')

# To adjust subplot parameters to give specified padding and prevent overlap of
↪labels
plt.tight_layout()
plt.show()

```



### 0.3.2 Univariate Analysis - Categorical Columns

The distribution of categorical columns will be plotted using barcharts.

```
[45]: df2[cat_columns]
```

```
[45]:
```

	gender	region	marital_status	bmi_category	smoking_status	\
0	Male	Northwest	Unmarried	Normal	No Smoking	
1	Female	Southeast	Married	Obesity	Regular	
2	Female	Northeast	Married	Normal	No Smoking	
3	Female	Southeast	Married	Normal	No Smoking	
4	Male	Northeast	Unmarried	Overweight	Regular	
...	...	...	...	...	...	

49903	Female	Northwest	Unmarried	Underweight	No Smoking
49904	Female	Southeast	Married	Normal	No Smoking
49905	Male	Northwest	Unmarried	Normal	Regular
49906	Male	Northwest	Unmarried	Normal	No Smoking
49907	Female	Southwest	Married	Normal	Occasional

	employment_status	income_level	medical_history	insurance_plan
0	Salaried	<10L	Diabetes	Bronze
1	Salaried	<10L	Diabetes	Bronze
2	Self-Employed	10L - 25L	High blood pressure	Silver
3	Salaried	> 40L	No Disease	Gold
4	Self-Employed	> 40L	High blood pressure	Silver
...	...	...	...	...
49903	Self-Employed	25L - 40L	No Disease	Bronze
49904	Salaried	> 40L	Thyroid	Gold
49905	Freelancer	25L - 40L	No Disease	Bronze
49906	Salaried	10L - 25L	No Disease	Bronze
49907	Self-Employed	<10L	Diabetes	Silver

[49908 rows x 9 columns]

In previous runs, the x-axis labels for the `medical_history` column were misaligned in the plots. To resolve this issue, the values will be mapped using the dictionary `ds_dict` as shown below.

```
[46]: ds_dict = {'No Disease': 'NO DISEASE',
'Diabetes': 'DB',
'High blood pressure': 'HBP',
'Thyroid': 'THY',
'Heart disease': 'HD',
'Diabetes & High blood pressure': 'DB & HBP',
'High blood pressure & Heart disease': 'HBP & HD',
'Diabetes & Thyroid': 'DB & THY',
'Diabetes & Heart disease': 'DB & HD'}
```

`ds_dict`

```
[46]: {'No Disease': 'NO DISEASE',
'Diabetes': 'DB',
'High blood pressure': 'HBP',
'Thyroid': 'THY',
'Heart disease': 'HD',
'Diabetes & High blood pressure': 'DB & HBP',
'High blood pressure & Heart disease': 'HBP & HD',
'Diabetes & Thyroid': 'DB & THY',
'Diabetes & Heart disease': 'DB & HD'}
```

```
[47]: # Showing how mapping works
```

```
print('Without Mapping:')
print(list(df2['medical_history'].value_counts(normalize=True).index))
print('\n','*' * 100, '\n')
print('With Mapping:')
print(list(df2['medical_history'].value_counts(normalize=True).index.
        ↪map(ds_dict)))
```

Without Mapping:

```
['No Disease', 'Diabetes', 'High blood pressure', 'Thyroid', 'Heart disease',
'Diabetes & High blood pressure', 'High blood pressure & Heart disease',
'Diabetes & Thyroid', 'Diabetes & Heart disease']
```

```
*****
*****
```

With Mapping:

```
['NO DISEASE', 'DB', 'HBP', 'THY', 'HD', 'DB & HBP', 'HBP & HD', 'DB & THY', 'DB
& HD']
```

```
[48]: # Plotting Distribution for all categorical columns
```

```
fig, ax = plt.subplots(3,3,figsize=(16,22))

# Iterator initiation to retrieve one value at a time
it = iter(cat_columns)

for i in range(3):
    for j in range(3):
        # To retrieve one value at a time
        col = next(it)

        # Format column name for display: capitalize first letter of each word
        ↪and replace underscores with spaces
        x_ax = col.title().replace('_', ' ')

        # Map the 'medical_history' column using ds_dict to ensure consistent
        ↪and clean labels for plotting

        if col == 'medical_history':
            x_val = df2[col].value_counts(normalize=True).index.map(ds_dict)
        else:
            x_val = df2[col].value_counts(normalize=True).index

        h_val = df2[col].value_counts(normalize=True).values * 100
        ax[i,j].bar(x=x_val,height=h_val)
```

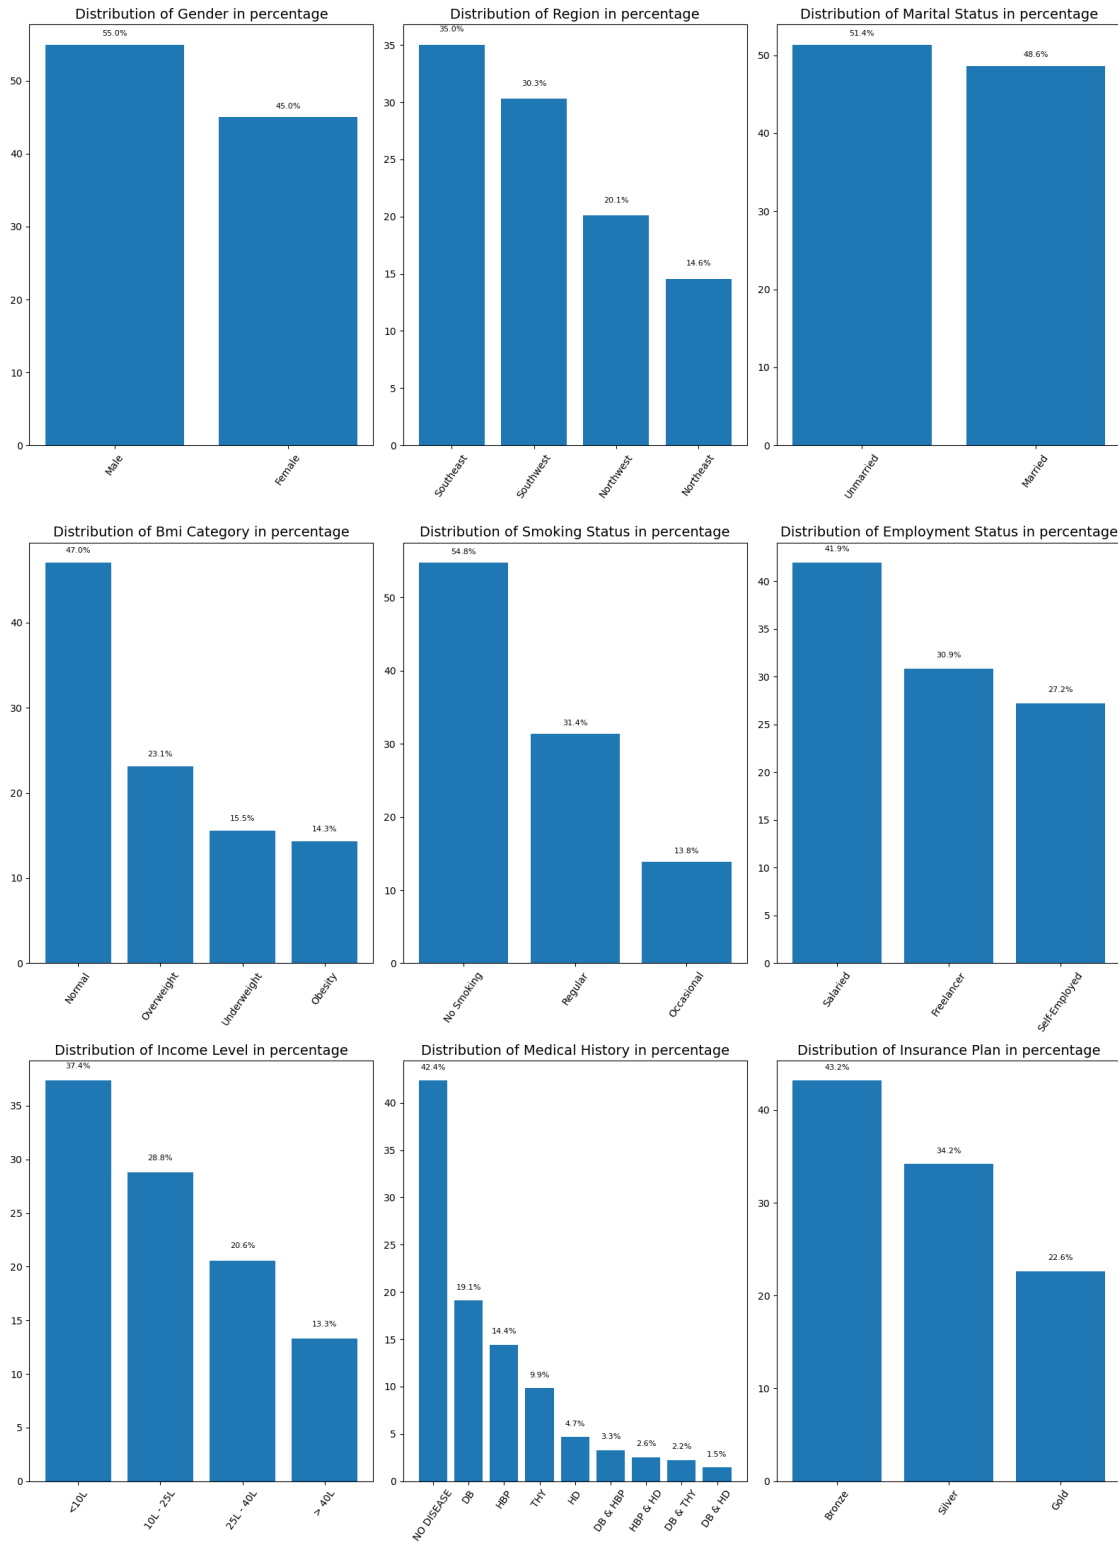
```

ax[i,j].set_title(f'Distribution of {x_ax} in percentage',fontsize=14)
ax[i,j].tick_params(axis='x', rotation=55)

# Annotate each bar with its corresponding percentage value for better
↳ interpretability
    for idx, val in enumerate(h_val):
        ax[i,j].text(idx, val + 1, f'{round(val,1)}%', ha='center',
        ↳va='bottom', fontsize=8)

plt.tight_layout()
plt.show()

```



### 0.3.3 Bivariate Analysis - Numeric Columns

The relationship between each numerical column (excluding the `annual_premium_amount` column itself) and the `annual_premium_amount` column will be plotted to analyze potential correlations.

```
[49]: df2[numeric_columns]
```

```
[49]:
```

	age	number_of_dependants	income_lakhs	annual_premium_amount
0	26	0	6	9053
1	29	2	6	16339
2	49	2	20	18164
3	30	3	77	20303
4	18	0	99	13365
...	...	...	...	...
49903	24	0	35	9111
49904	47	2	82	27076
49905	21	0	32	8564
49906	18	2	20	9490
49907	48	3	7	19730

```
[49908 rows x 4 columns]
```

```
[50]: # Plotting the scatter plots

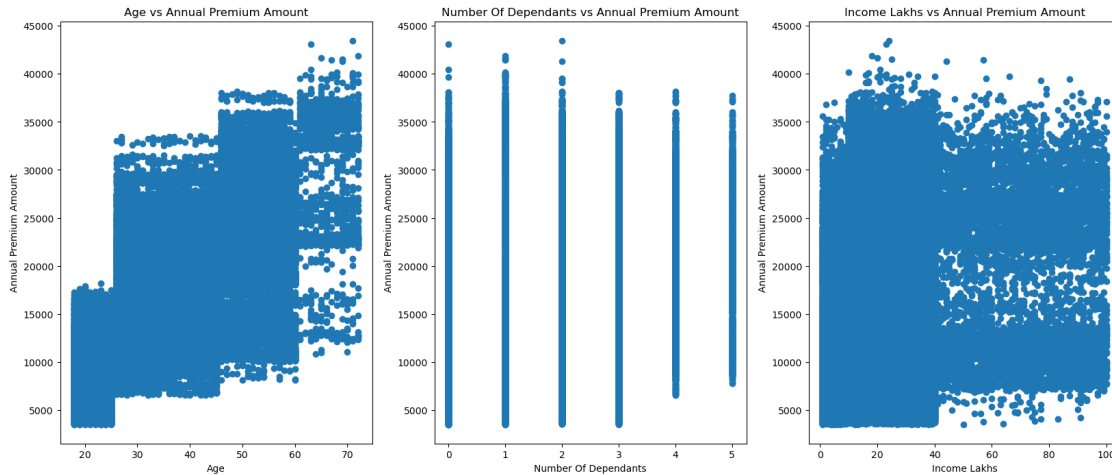
fig, ax = plt.subplots(1,3,figsize=(20,8))

for i in range(3):
    col = numeric_columns[i]

    # Format column names for display: capitalize first letter of each word and
    # replace underscores with spaces
    x_ax = col.title().replace('_', ' ')
    y_ax = 'annual_premium_amount'.title().replace('_', ' ')

    ax[i].scatter(x=df2[col],y=df2['annual_premium_amount'])
    ax[i].set_title(f'{x_ax} vs {y_ax}')
    ax[i].set_xlabel(x_ax)
    ax[i].set_ylabel(y_ax)
```





### 0.3.4 Bivariate Analysis - Categorical

The relationship between the categorical columns `income_level` and `insurance_plan` will be explored using multiple approaches:

- Cross-tabulation
- Grouped Bar Chart
- Stacked Bar Chart
- Heatmap

```
[51]: df2.head()
```

```
[51]:   age  gender  region marital_status  number_of_dependants  bmi_category \
0   26   Male  Northwest      Unmarried                0      Normal
1   29  Female  Southeast      Married                2      Obesity
2   49  Female  Northeast      Married                2      Normal
3   30  Female  Southeast      Married                3      Normal
4   18   Male  Northeast      Unmarried                0  Overweight
```

```
   smoking_status  employment_status  income_level  income_lakhs  \
0   No Smoking      Salaried      <10L             6
1   Regular      Salaried      <10L             6
2   No Smoking  Self-Employed  10L - 25L            20
3   No Smoking      Salaried      > 40L            77
4   Regular  Self-Employed      > 40L            99
```

```
   medical_history  insurance_plan  annual_premium_amount
0      Diabetes      Bronze      9053
1      Diabetes      Bronze     16339
2  High blood pressure      Silver     18164
3      No Disease      Gold      20303
```

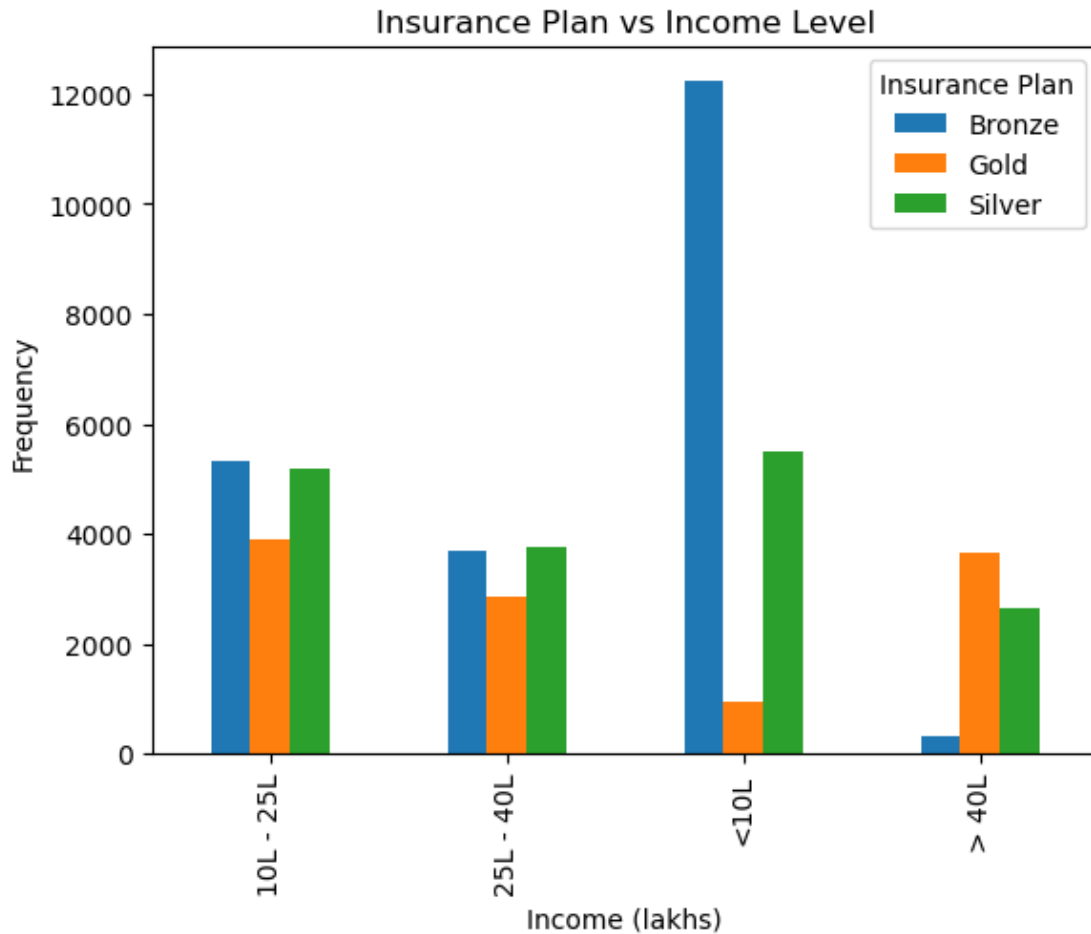
**Cross-tabulation** The frequency distribution between `income_level` and `insurance_plan` will be examined using a cross-tabulation, which will display the number of individuals in each combination of income level and insurance plan.

```
[52]: ct = pd.crosstab(df['income_level'],df['insurance_plan'])
      ct
```

```
[52]: insurance_plan  Bronze  Gold  Silver
income_level
10L - 25L           5314  3886    5189
25L - 40L           3686  2844    3753
<10L                12239   931    5495
> 40L                330   3662    2647
```

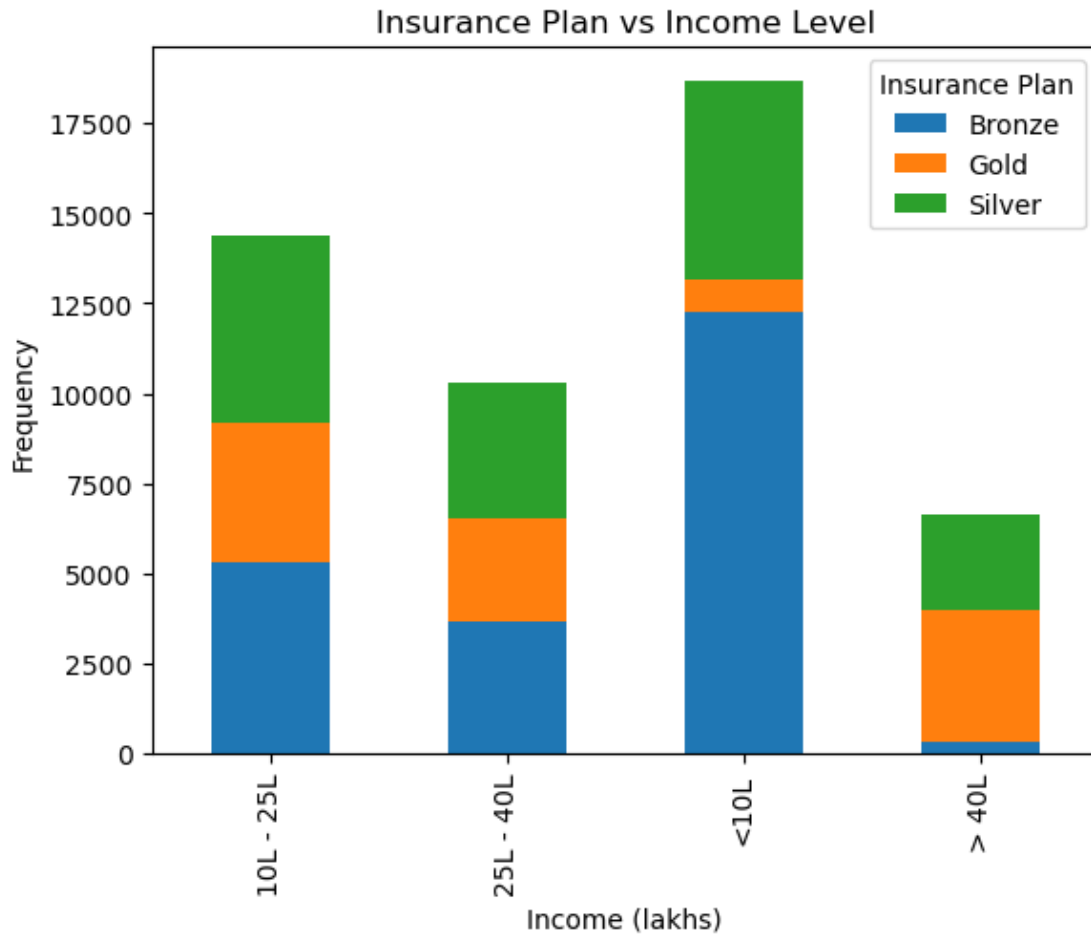
**Grouped Bar Chart** The count of each `insurance_plan` across different `income_level` categories will be visualized using grouped bar charts, where each income level will have side-by-side bars representing the different insurance plans for easy comparison.

```
[53]: ct.plot(kind='bar')
      plt.title('Insurance Plan vs Income Level')
      plt.xlabel('Income (lakhs)')
      plt.ylabel('Frequency')
      plt.legend(title='Insurance Plan')
      plt.show()
```




**Stacked Bar Chart** The distribution of `insurance_plan` within each `income_level` will be represented in a stacked bar chart format, helping to understand the proportion of each plan type within the income categories.

```
[54]: ct.plot(kind='bar',stacked=True)
plt.title('Insurance Plan vs Income Level')
plt.xlabel('Income (lakhs)')
plt.ylabel('Frequency')
plt.legend(title='Insurance Plan')
plt.show()
```



**Heatmap** The intensity of the relationship between `insurance_plan` and `income_level` will be visualized using a heatmap, where the cells are color-coded based on frequency to highlight patterns and concentrations in the data.

```
[55]: sns.heatmap(ct,annot=True,fmt='0')
plt.title('Insurance Plan vs Income Level')
plt.xlabel('Insurance Plan')
plt.ylabel('Income Level')
plt.show()
```



## 0.4 Feature Engineering

The following feature engineering steps were performed:

- Created new features as required
- Transformed features, including encoding categorical variables
- Selected important features using statistical methods such as Variance Inflation Factor (VIF) and correlation analysis

### 0.4.1 Feature Creation - 'total\_risk\_score'

In this we created a new feature `total_risk_score` from `medical_history` column

```
[56]: # Extract all distinct medical conditions listed in the dataset
```

```
df2['medical_history'].unique()
```

```
[56]: array(['Diabetes', 'High blood pressure', 'No Disease',
        'Diabetes & High blood pressure', 'Thyroid', 'Heart disease',
        'High blood pressure & Heart disease', 'Diabetes & Thyroid',
        'Diabetes & Heart disease'], dtype=object)
```

The following risk scores, as provided by the business, will be assigned to the corresponding medical conditions:

- Heart disease: 8
- Diabetes: 6
- High blood pressure: 6
- Thyroid: 5
- No Disease: 0
- None: 0

```
[57]: df3 = df2.copy()
df3.head()
```

```
[57]:   age  gender  region marital_status  number_of_dependants  bmi_category \
0   26   Male  Northwest      Unmarried                0      Normal
1   29  Female  Southeast      Married                2      Obesity
2   49  Female  Northeast      Married                2      Normal
3   30  Female  Southeast      Married                3      Normal
4   18   Male  Northeast      Unmarried                0  Overweight

   smoking_status  employment_status  income_level  income_lakhs  \
0      No Smoking      Salaried      <10L                6
1      Regular      Salaried      <10L                6
2      No Smoking  Self-Employed    10L - 25L            20
3      No Smoking      Salaried      > 40L            77
4      Regular  Self-Employed      > 40L            99

   medical_history  insurance_plan  annual_premium_amount
0          Diabetes      Bronze          9053
1          Diabetes      Bronze          16339
2  High blood pressure      Silver          18164
3          No Disease      Gold          20303
4  High blood pressure      Silver          13365
```

```
[58]: # Split the 'medical_history' column into 'disease1' and 'disease2' using '&'
      ↪ as the delimiter

df3[['disease1','disease2']] = df3['medical_history'].str.lower().str.split(' &
      ↪ ',expand=True)
df3.head()
```

```
[58]:   age  gender  region marital_status  number_of_dependants  bmi_category \
0   26   Male  Northwest      Unmarried                0      Normal
1   29  Female  Southeast      Married                2      Obesity
2   49  Female  Northeast      Married                2      Normal
3   30  Female  Southeast      Married                3      Normal
4   18   Male  Northeast      Unmarried                0  Overweight
```

	smoking_status	employment_status	income_level	income_lakhs \
0	No Smoking	Salaried	<10L	6
1	Regular	Salaried	<10L	6
2	No Smoking	Self-Employed	10L - 25L	20
3	No Smoking	Salaried	> 40L	77
4	Regular	Self-Employed	> 40L	99

	medical_history	insurance_plan	annual_premium_amount \
0	Diabetes	Bronze	9053
1	Diabetes	Bronze	16339
2	High blood pressure	Silver	18164
3	No Disease	Gold	20303
4	High blood pressure	Silver	13365

	disease1	disease2
0	diabetes	None
1	diabetes	None
2	high blood pressure	None
3	no disease	None
4	high blood pressure	None

[59]: *# Risk Score Dictionary*

```

risk_score_dict = {
    'heart disease' : 8,
    'diabetes' : 6,
    'high blood pressure' : 6,
    'thyroid' : 5,
    'no disease' : 0,
    None : 0
}

```

[60]: *# Mapping each disease to its corresponding score using a predefined dictionary*

```

df3['disease1_score'] = df3['disease1'].map(risk_score_dict)
df3['disease2_score'] = df3['disease2'].map(risk_score_dict)
df3.sample(2)

```

[60]:

	age	gender	region	marital_status	number_of_dependants \
11647	53	Male	Southeast	Married	3
5418	59	Male	Southeast	Married	2

	bmi_category	smoking_status	employment_status	income_level \
11647	Obesity	Occasional	Salaried	<10L
5418	Normal	Regular	Self-Employed	<10L

	income_lakhs	medical_history	insurance_plan	annual_premium_amount \
--	--------------	-----------------	----------------	-------------------------

11647	3	Thyroid	Bronze	15240
5418	3	Thyroid	Bronze	14013

	disease1	disease2	disease1_score	disease2_score
11647	thyroid	None	5	0
5418	thyroid	None	5	0

```
[61]: # Check if all diseases have been assigned a score and identify any missing
      ↪ values
```

```
print('Unique Scores in Disease1 -> ',df3['disease1_score'].unique())
print('Unique Scores in Disease2 -> ',df3['disease2_score'].unique())
```

```
Unique Scores in Disease1 -> [6 0 5 8]
Unique Scores in Disease2 -> [0 6 8 5]
```

Since there are no NaN values, it can be concluded that all entries in the `disease1` and `disease2` columns have been successfully mapped.

```
[62]: # Calculating the total risk score by summing 'disease1_score' and
      ↪ 'disease2_score'

df3['total_risk_score'] = df3['disease1_score'] + df3['disease2_score']
df3.sample(4)
```

```
[62]:      age gender      region marital_status  number_of_dependants \
2358    69   Male  Southeast      Married                1
1268    60   Male  Southwest      Married                3
9998    22   Male  Southwest      Married                3
27335   56   Male  Southwest      Married                5
```

	bmi_category	smoking_status	employment_status	income_level	\
2358	Normal	No Smoking	Self-Employed	25L - 40L	
1268	Normal	No Smoking	Self-Employed	> 40L	
9998	Normal	No Smoking	Freelancer	<10L	
27335	Normal	Regular	Salaried	<10L	

	income_lakhs	medical_history	insurance_plan	\
2358	28	Thyroid	Bronze	
1268	45	High blood pressure	Gold	
9998	7	Diabetes	Bronze	
27335	8	High blood pressure	Silver	

	annual_premium_amount	disease1	disease2	disease1_score	\
2358	12709	thyroid	None	5	
1268	26541	high blood pressure	None	6	
9998	9001	diabetes	None	6	
27335	21654	high blood pressure	None	6	



	disease2_score	total_risk_score
2358	0	5
1268	0	6
9998	0	6
27335	0	6

The following columns are being dropped as they are no longer needed: `medical_history`, `disease1`, `disease2`, `disease1_score`, and `disease2_score`.

```
[63]: cols_to_drop = ['medical_history', 'disease1', 'disease2', 'disease1_score',
↳ 'disease2_score']
df4 = df3.drop(cols_to_drop,axis=1)
df4.reset_index(drop=True,inplace=True)
df4
```

```
[63]:
```

	age	gender	region	marital_status	number_of_dependants	\
0	26	Male	Northwest	Unmarried		0
1	29	Female	Southeast	Married		2
2	49	Female	Northeast	Married		2
3	30	Female	Southeast	Married		3
4	18	Male	Northeast	Unmarried		0
...	...	...	...	...	...	
49903	24	Female	Northwest	Unmarried		0
49904	47	Female	Southeast	Married		2
49905	21	Male	Northwest	Unmarried		0
49906	18	Male	Northwest	Unmarried		2
49907	48	Female	Southwest	Married		3

	bmi_category	smoking_status	employment_status	income_level	\
0	Normal	No Smoking	Salaried	<10L	
1	Obesity	Regular	Salaried	<10L	
2	Normal	No Smoking	Self-Employed	10L - 25L	
3	Normal	No Smoking	Salaried	> 40L	
4	Overweight	Regular	Self-Employed	> 40L	
...	...	...	...	...	
49903	Underweight	No Smoking	Self-Employed	25L - 40L	
49904	Normal	No Smoking	Salaried	> 40L	
49905	Normal	Regular	Freelancer	25L - 40L	
49906	Normal	No Smoking	Salaried	10L - 25L	
49907	Normal	Occasional	Self-Employed	<10L	

	income_lakhs	insurance_plan	annual_premium_amount	total_risk_score
0	6	Bronze	9053	6
1	6	Bronze	16339	6
2	20	Silver	18164	6
3	77	Gold	20303	0
4	99	Silver	13365	6

...	...	...	...	...
49903	35	Bronze	9111	0
49904	82	Gold	27076	5
49905	32	Bronze	8564	0
49906	20	Bronze	9490	0
49907	7	Silver	19730	6

[49908 rows x 13 columns]

## 0.4.2 Feature Cleaning & Transformation

The following transformations were applied to the dataset features:

- Label Encoding for ordinal categorical variables (to preserve order)
- One Hot Encoding for nominal categorical variables (to avoid introducing ordinal relationships)

```
[64]: df5 = df4.copy()
df5.head()
```

```
[64]:   age  gender  region marital_status  number_of_dependants bmi_category \
0   26   Male  Northwest      Unmarried                0      Normal
1   29  Female  Southeast        Married                2      Obesity
2   49  Female  Northeast        Married                2      Normal
3   30  Female  Southeast        Married                3      Normal
4   18   Male  Northeast      Unmarried                0  Overweight

   smoking_status employment_status income_level  income_lakhs insurance_plan \
0      No Smoking          Salaried      <10L             6      Bronze
1      Regular          Salaried      <10L             6      Bronze
2      No Smoking      Self-Employed    10L - 25L          20      Silver
3      No Smoking          Salaried      > 40L          77      Gold
4      Regular      Self-Employed      > 40L          99      Silver

   annual_premium_amount  total_risk_score
0              9053             6
1             16339             6
2             18164             6
3             20303             0
4             13365             6
```

### Label Encodig - 'income\_level'

```
[65]: # Extract all distinct income levels listed in the dataset

df5.income_level.unique()
```

```
[65]: array(['<10L', '10L - 25L', '> 40L', '25L - 40L'], dtype=object)
```

```
[66]: # Income level dictionary
```

```
income_level_dict = {  
    '<10L' : 1,  
    '10L - 25L' : 2,  
    '25L - 40L' : 3,  
    '> 40L' : 4  
}
```

```
[67]: # Mapping each income level to a value using a predefined dictionary
```

```
df5['income_level'] = df5['income_level'].map(income_level_dict)
```

```
[68]: # After mapping
```

```
df5.income_level.unique()
```

```
[68]: array([1, 2, 4, 3])
```

### Label Encodig - 'insurance\_plan'

```
[69]: # Extract all distinct insurance plan listed in the dataset
```

```
df5.insurance_plan.unique()
```

```
[69]: array(['Bronze', 'Silver', 'Gold'], dtype=object)
```

```
[70]: # Insurance Plan dictionary
```

```
insurance_plan_dict = {  
    'Bronze' : 1,  
    'Silver' : 2,  
    'Gold' : 3,  
}
```

```
[71]: # Mapping each insurance plan to a value using a predefined dictionary
```

```
df5['insurance_plan'] = df5['insurance_plan'].map(insurance_plan_dict)
```

```
[72]: # After mapping
```

```
df5.insurance_plan.unique()
```

```
[72]: array([1, 2, 3])
```

### One Hot Encoding

```
[73]: # Selecting columns to perform one hot encoding
```

```
cols_to_encode = ['gender', 'region', 'marital_status', 'bmi_category',
                  'smoking_status', 'employment_status']
cols_to_encode
```

```
[73]: ['gender',
       'region',
       'marital_status',
       'bmi_category',
       'smoking_status',
       'employment_status']
```

```
[74]: # Performing One Hot Encoding on df5
```

```
df6 = pd.get_dummies(df5, columns = cols_to_encode, dtype=int, drop_first=True)
```

```
[75]: df6.sample(5)
```

```
[75]:
```

	age	number_of_dependants	income_level	income_lakhs	insurance_plan	\
41672	41	1	4	72	3	
1380	20	0	3	27	1	
9953	18	2	1	7	1	
7862	22	2	3	31	1	
16612	53	2	2	10	3	

	annual_premium_amount	total_risk_score	gender_Male	region_Northwest	\
41672	21674	0	0	0	
1380	5773	0	1	0	
9953	9530	0	1	0	
7862	9690	0	0	0	
16612	29848	6	1	0	

	region_Southeast	region_Southwest	marital_status_Unmarried	\
41672	0	1	1	
1380	0	1	1	
9953	1	0	1	
7862	1	0	0	
16612	1	0	0	

	bmi_category_Obesity	bmi_category_Overweight	\
41672	0	0	
1380	0	0	
9953	0	0	
7862	0	1	
16612	0	0	

	bmi_category_Underweight	smoking_status_Occasional	\
41672	1	1	

1380	1	0
9953	1	0
7862	0	0
16612	0	0

	smoking_status_Regular	employment_status_Salaried \
41672	0	0
1380	1	1
9953	0	1
7862	0	0
16612	1	0

	employment_status_Self-Employed
41672	1
1380	0
9953	0
7862	1
16612	1

[76]: df6.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49908 entries, 0 to 49907
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                  49908 non-null  int64
1   number_of_dependants                 49908 non-null  int64
2   income_level                         49908 non-null  int64
3   income_lakhs                        49908 non-null  int64
4   insurance_plan                      49908 non-null  int64
5   annual_premium_amount               49908 non-null  int64
6   total_risk_score                    49908 non-null  int64
7   gender_Male                         49908 non-null  int64
8   region_Northwest                    49908 non-null  int64
9   region_Southeast                    49908 non-null  int64
10  region_Southwest                    49908 non-null  int64
11  marital_status_Unmarried            49908 non-null  int64
12  bmi_category_Obesity                49908 non-null  int64
13  bmi_category_Overweight             49908 non-null  int64
14  bmi_category_Underweight            49908 non-null  int64
15  smoking_status_Occasional           49908 non-null  int64
16  smoking_status_Regular              49908 non-null  int64
17  employment_status_Salaried          49908 non-null  int64
18  employment_status_Self-Employed    49908 non-null  int64
dtypes: int64(19)
memory usage: 7.2 MB
```

```
[77]: # Showing all the non-encoded columns and one encoded columns from each category

sampled_encoded_cols_index = [0,1,2,3,4,5,6,7,8,11,12,15,17]

df6.iloc[0:5,sampled_encoded_cols_index]
```

```
[77]:  age  number_of_dependants  income_level  income_lakhs  insurance_plan  \
0    26                      0              1              6              1
1    29                      2              1              6              1
2    49                      2              2             20              2
3    30                      3              4             77              3
4    18                      0              4             99              2

    annual_premium_amount  total_risk_score  gender_Male  region_Northwest  \
0                9053              6              1              1
1               16339              6              0              0
2               18164              6              0              0
3               20303              0              0              0
4               13365              6              1              0

    marital_status_Unmarried  bmi_category_Obesity  smoking_status_Occasional  \
0                        1                      0                        0
1                        0                      1                        0
2                        0                      0                        0
3                        0                      0                        0
4                        1                      0                        0

    employment_status_Salaried
0                        1
1                        1
2                        0
3                        1
4                        0
```

### 0.4.3 Feature Selection

- To identify the most relevant features, both pairwise correlations and multicollinearity will be analyzed.
- *Correlation* analysis will be used to detect linear relationships, while *Variance Inflation Factor (VIF)* will be employed to identify multicollinearity.
- Prior to these analyses, features will be scaled to ensure comparability across variables.

```
[78]: # Before Scaling the features

df7 = df6.copy()
df7.sample(3)
```

```

[78]:      age  number_of_dependants  income_level  income_lakhs  insurance_plan  \
22498    21                      3              3             28           1
6982     20                      1              3             27           1
15687    50                      3              3             26           3

      annual_premium_amount  total_risk_score  gender_Male  region_Northwest  \
22498                   8601                0            1                0
6982                   7292                0            1                0
15687                  28031                6            1                0

      region_Southeast  region_Southwest  marital_status_Unmarried  \
22498                0                1                        0
6982                0                1                        1
15687                1                0                        0

      bmi_category_Obesity  bmi_category_Overweight  \
22498                    0                        0
6982                    0                        0
15687                    0                        0

      bmi_category_Underweight  smoking_status_Occasional  \
22498                        0                        0
6982                        0                        0
15687                        1                        1

      smoking_status_Regular  employment_status_Salaried  \
22498                      0                        0
6982                      1                        0
15687                      0                        1

      employment_status_Self-Employed
22498                                0
6982                                1
15687                                0

```

```
[79]: df7.columns
```

```

[79]: Index(['age', 'number_of_dependants', 'income_level', 'income_lakhs',
          'insurance_plan', 'annual_premium_amount', 'total_risk_score',
          'gender_Male', 'region_Northwest', 'region_Southeast',
          'region_Southwest', 'marital_status_Unmarried', 'bmi_category_Obesity',
          'bmi_category_Overweight', 'bmi_category_Underweight',
          'smoking_status_Occasional', 'smoking_status_Regular',
          'employment_status_Salaried', 'employment_status_Self-Employed'],
          dtype='object')

```

The following columns will be scaled: age, number\_of\_dependants, income\_level, income\_lakhs, insurance\_plan, total\_risk\_score.

Scaling type : Minmax scaling

```
[80]: # Scaling using 'MinMaxScaler'

cols_to_scale = ['age', 'number_of_dependants', 'income_level', 'income_lakhs',
                 'insurance_plan', 'total_risk_score']
mms = MinMaxScaler()
df7[cols_to_scale] = mms.fit_transform(df7[cols_to_scale])
```

## Correlation

```
[81]: # Correlation Matrix between features

cr = df7.corr()
cr
```

```
[81]:
```

	age	number_of_dependants	income_level	\
age	1.000000	0.415742	0.029851	
number_of_dependants	0.415742	1.000000	0.006564	
income_level	0.029851	0.006564	1.000000	
income_lakhs	0.025060	0.006074	0.906830	
insurance_plan	0.496317	0.256459	0.440428	
annual_premium_amount	0.767569	0.414691	0.271811	
total_risk_score	0.442773	0.371498	0.013506	
gender_Male	-0.002219	-0.003093	0.063108	
region_Northwest	0.000464	0.001693	-0.003324	
region_Southeast	0.003305	0.003620	-0.000259	
region_Southwest	-0.003424	-0.000339	0.009367	
marital_status_Unmarried	-0.543104	-0.841717	-0.012994	
bmi_category_Obesity	0.152496	0.115397	-0.002244	
bmi_category_Overweight	0.153148	0.110451	0.007947	
bmi_category_Underweight	-0.115888	-0.093881	0.000350	
smoking_status_Occasional	0.066596	0.071762	-0.001340	
smoking_status_Regular	0.059380	0.094829	0.020275	
employment_status_Salaried	-0.008093	0.067066	-0.134032	
employment_status_Self-Employed	0.314684	0.115930	0.139333	

	income_lakhs	insurance_plan	\
age	0.025060	0.496317	
number_of_dependants	0.006074	0.256459	
income_level	0.906830	0.440428	
income_lakhs	1.000000	0.410753	
insurance_plan	0.410753	1.000000	
annual_premium_amount	0.243058	0.834148	
total_risk_score	0.009626	0.260932	
gender_Male	0.039126	0.034211	
region_Northwest	-0.005192	-0.002821	
region_Southeast	-0.001250	0.004082	



region_Southwest	0.009929	-0.000977
marital_status_Unmarried	-0.011099	-0.316800
bmi_category_Obesity	0.000314	0.094698
bmi_category_Overweight	0.007150	0.098639
bmi_category_Underweight	-0.000740	-0.073881
smoking_status_Occasional	0.002306	0.037351
smoking_status_Regular	0.010948	0.059587
employment_status_Salaried	-0.100510	-0.041582
employment_status_Self-Employed	0.109759	0.223947

	annual_premium_amount	total_risk_score \
age	0.767569	0.442773
number_of_dependants	0.414691	0.371498
income_level	0.271811	0.013506
income_lakhs	0.243058	0.009626
insurance_plan	0.834148	0.260932
annual_premium_amount	1.000000	0.519458
total_risk_score	0.519458	1.000000
gender_Male	0.064470	-0.003754
region_Northwest	-0.005078	-0.005627
region_Southeast	0.008235	0.002019
region_Southwest	-0.003828	-0.000132
marital_status_Unmarried	-0.516350	-0.433916
bmi_category_Obesity	0.249847	0.101039
bmi_category_Overweight	0.187103	0.102556
bmi_category_Underweight	-0.135289	-0.087996
smoking_status_Occasional	0.060610	0.071090
smoking_status_Regular	0.198829	0.093822
employment_status_Salaried	-0.005442	0.059511
employment_status_Self-Employed	0.289438	0.135824

	gender_Male	region_Northwest \
age	-0.002219	0.000464
number_of_dependants	-0.003093	0.001693
income_level	0.063108	-0.003324
income_lakhs	0.039126	-0.005192
insurance_plan	0.034211	-0.002821
annual_premium_amount	0.064470	-0.005078
total_risk_score	-0.003754	-0.005627
gender_Male	1.000000	-0.003287
region_Northwest	-0.003287	1.000000
region_Southeast	-0.002988	-0.368277
region_Southwest	0.001873	-0.330654
marital_status_Unmarried	0.003944	-0.001083
bmi_category_Obesity	-0.031351	-0.002262
bmi_category_Overweight	0.080588	0.001194
bmi_category_Underweight	-0.043000	-0.002017

smoking_status_Occasional	-0.045618	-0.000669
smoking_status_Regular	0.305180	-0.000255
employment_status_Salaried	0.005559	0.004574
employment_status_Self-Employed	0.001055	0.000110

	region_Southeast	region_Southwest	\
age	0.003305	-0.003424	
number_of_dependants	0.003620	-0.000339	
income_level	-0.000259	0.009367	
income_lakhs	-0.001250	0.009929	
insurance_plan	0.004082	-0.000977	
annual_premium_amount	0.008235	-0.003828	
total_risk_score	0.002019	-0.000132	
gender_Male	-0.002988	0.001873	
region_Northwest	-0.368277	-0.330654	
region_Southeast	1.000000	-0.484271	
region_Southwest	-0.484271	1.000000	
marital_status_Unmarried	-0.003980	0.004287	
bmi_category_Obesity	0.006373	-0.000934	
bmi_category_Overweight	0.002511	-0.002208	
bmi_category_Underweight	0.000211	-0.000499	
smoking_status_Occasional	-0.002191	0.006282	
smoking_status_Regular	-0.000401	-0.001858	
employment_status_Salaried	-0.006575	0.000249	
employment_status_Self-Employed	0.003287	-0.005618	

	marital_status_Unmarried	\
age	-0.543104	
number_of_dependants	-0.841717	
income_level	-0.012994	
income_lakhs	-0.011099	
insurance_plan	-0.316800	
annual_premium_amount	-0.516350	
total_risk_score	-0.433916	
gender_Male	0.003944	
region_Northwest	-0.001083	
region_Southeast	-0.003980	
region_Southwest	0.004287	
marital_status_Unmarried	1.000000	
bmi_category_Obesity	-0.118092	
bmi_category_Overweight	-0.117312	
bmi_category_Underweight	0.089315	
smoking_status_Occasional	-0.075253	
smoking_status_Regular	-0.092448	
employment_status_Salaried	-0.055285	
employment_status_Self-Employed	-0.171646	

	bmi_category_Obesity \
age	0.152496
number_of_dependants	0.115397
income_level	-0.002244
income_lakhs	0.000314
insurance_plan	0.094698
annual_premium_amount	0.249847
total_risk_score	0.101039
gender_Male	-0.031351
region_Northwest	-0.002262
region_Southeast	0.006373
region_Southwest	-0.000934
marital_status_Unmarried	-0.118092
bmi_category_Obesity	1.000000
bmi_category_Overweight	-0.224205
bmi_category_Underweight	-0.175299
smoking_status_Occasional	0.028321
smoking_status_Regular	0.017526
employment_status_Salaried	0.021085
employment_status_Self-Employed	0.040750

	bmi_category_Overweight \
age	0.153148
number_of_dependants	0.110451
income_level	0.007947
income_lakhs	0.007150
insurance_plan	0.098639
annual_premium_amount	0.187103
total_risk_score	0.102556
gender_Male	0.080588
region_Northwest	0.001194
region_Southeast	0.002511
region_Southwest	-0.002208
marital_status_Unmarried	-0.117312
bmi_category_Obesity	-0.224205
bmi_category_Overweight	1.000000
bmi_category_Underweight	-0.235191
smoking_status_Occasional	0.029486
smoking_status_Regular	0.070944
employment_status_Salaried	0.028753
employment_status_Self-Employed	0.035602

	bmi_category_Underweight \
age	-0.115888
number_of_dependants	-0.093881
income_level	0.000350
income_lakhs	-0.000740

insurance_plan	-0.073881
annual_premium_amount	-0.135289
total_risk_score	-0.087996
gender_Male	-0.043000
region_Northwest	-0.002017
region_Southeast	0.000211
region_Southwest	-0.000499
marital_status_Unmarried	0.089315
bmi_category_Obesity	-0.175299
bmi_category_Overweight	-0.235191
bmi_category_Underweight	1.000000
smoking_status_Occasional	-0.023061
smoking_status_Regular	-0.044704
employment_status_Salaried	-0.023124
employment_status_Self-Employed	-0.023601

	smoking_status_Occasional \
age	0.066596
number_of_dependants	0.071762
income_level	-0.001340
income_lakhs	0.002306
insurance_plan	0.037351
annual_premium_amount	0.060610
total_risk_score	0.071090
gender_Male	-0.045618
region_Northwest	-0.000669
region_Southeast	-0.002191
region_Southwest	0.006282
marital_status_Unmarried	-0.075253
bmi_category_Obesity	0.028321
bmi_category_Overweight	0.029486
bmi_category_Underweight	-0.023061
smoking_status_Occasional	1.000000
smoking_status_Regular	-0.270923
employment_status_Salaried	0.026424
employment_status_Self-Employed	0.012346

	smoking_status_Regular \
age	0.059380
number_of_dependants	0.094829
income_level	0.020275
income_lakhs	0.010948
insurance_plan	0.059587
annual_premium_amount	0.198829
total_risk_score	0.093822
gender_Male	0.305180
region_Northwest	-0.000255

region_Southeast	-0.000401
region_Southwest	-0.001858
marital_status_Unmarried	-0.092448
bmi_category_Obesity	0.017526
bmi_category_Overweight	0.070944
bmi_category_Underweight	-0.044704
smoking_status_Occasional	-0.270923
smoking_status_Regular	1.000000
employment_status_Salaried	0.042991
employment_status_Self-Employed	0.002693

	employment_status_Salaried \
age	-0.008093
number_of_dependants	0.067066
income_level	-0.134032
income_lakhs	-0.100510
insurance_plan	-0.041582
annual_premium_amount	-0.005442
total_risk_score	0.059511
gender_Male	0.005559
region_Northwest	0.004574
region_Southeast	-0.006575
region_Southwest	0.000249
marital_status_Unmarried	-0.055285
bmi_category_Obesity	0.021085
bmi_category_Overweight	0.028753
bmi_category_Underweight	-0.023124
smoking_status_Occasional	0.026424
smoking_status_Regular	0.042991
employment_status_Salaried	1.000000
employment_status_Self-Employed	-0.519576

	employment_status_Self-Employed
age	0.314684
number_of_dependants	0.115930
income_level	0.139333
income_lakhs	0.109759
insurance_plan	0.223947
annual_premium_amount	0.289438
total_risk_score	0.135824
gender_Male	0.001055
region_Northwest	0.000110
region_Southeast	0.003287
region_Southwest	-0.005618
marital_status_Unmarried	-0.171646
bmi_category_Obesity	0.040750
bmi_category_Overweight	0.035602

```

bmi_category_Underweight          -0.023601
smoking_status_Occasional          0.012346
smoking_status_Regular             0.002693
employment_status_Salaried        -0.519576
employment_status_Self-Employed    1.000000

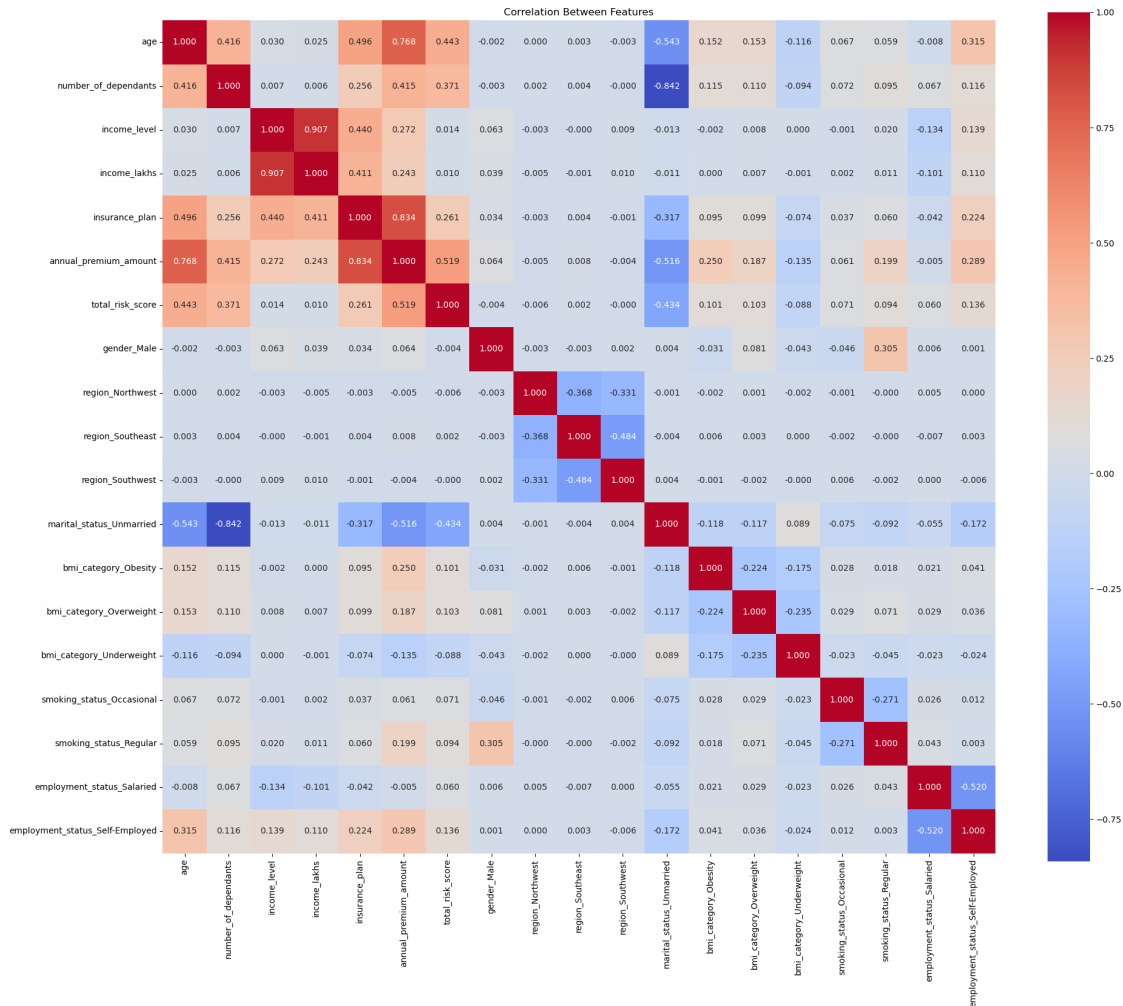
```

[82]: # Correlation Matrix displayed as a Heatmap

```

plt.figure(figsize=(20,20))
sns.heatmap(cr, annot=True, fmt='.3f', cmap='coolwarm', square=True,
            cbar_kws={"shrink": 0.75})
plt.title('Correlation Between Features')
plt.tight_layout()
plt.show()

```



```
[83]: # Display features that have a high correlation (|correlation| > 0.35) with
      ↪ 'annual_premium_amount'

      cr[abs(cr['annual_premium_amount']) > 0.35]['annual_premium_amount']
```

```
[83]: age                0.767569
      number_of_dependants  0.414691
      insurance_plan       0.834148
      annual_premium_amount 1.000000
      total_risk_score      0.519458
      marital_status_Unmarried -0.516350
      Name: annual_premium_amount, dtype: float64
```

```
[84]: # Extract the feature names (index labels) that have a high correlation
      ↪ (|correlation| > 0.35) with 'annual_premium_amount'

      cr[abs(cr['annual_premium_amount']) > 0.35]['annual_premium_amount'].index
```

```
[84]: Index(['age', 'number_of_dependants', 'insurance_plan',
          'annual_premium_amount', 'total_risk_score',
          'marital_status_Unmarried'],
          dtype='object')
```

### Observation

Based on the correlation heatmap, a strong positive correlation with `annual_premium_amount` is observed for the features `age`, `number_of_dependants`, `insurance_plan`, and `total_risk_score`, suggesting that these variables may significantly influence premium prediction.

Additionally, a strong negative correlation is shown by `marital_status_Unmarried`, indicating that being unmarried is generally associated with lower premium amounts.

**Variance Inflation Factor (VIF)** The Variance Inflation Factor (VIF) will be checked to detect and address multicollinearity among the features.

```
[85]: df7.columns
```

```
[85]: Index(['age', 'number_of_dependants', 'income_level', 'income_lakhs',
          'insurance_plan', 'annual_premium_amount', 'total_risk_score',
          'gender_Male', 'region_Northwest', 'region_Southeast',
          'region_Southwest', 'marital_status_Unmarried', 'bmi_category_Obesity',
          'bmi_category_Overweight', 'bmi_category_Underweight',
          'smoking_status_Occasional', 'smoking_status_Regular',
          'employment_status_Salaried', 'employment_status_Self-Employed'],
          dtype='object')
```

```
[86]: # Calculate VIF for each feature and store the results in a new DataFrame
```

```

# Initialize a dictionary to store feature names and their corresponding VIF
↪scores
vif_dict = {'features':[], 'vif_score':[]}

# Exclude the target variable from VIF calculation
temp_df = df7.drop('annual_premium_amount',axis=1)

# Loop through each feature to compute VIF
for i,col in enumerate(temp_df.columns):
    # Calculate the Variance Inflation Factor for the current feature
    vif = variance_inflation_factor(temp_df,i)

    # Append the feature name and its VIF score to the dictionary
    vif_dict['features'].append(col)
    vif_dict['vif_score'].append(vif)

```

```

[87]: # Convert the VIF dictionary into a DataFrame for better readability and
↪analysis
vif_df = pd.DataFrame(vif_dict)

# Sort features by their VIF scores in descending order to identify highly
↪collinear features
vif_df.sort_values(by='vif_score',ascending=False)

```

```

[87]:
           features  vif_score
2      income_level  12.450675
3      income_lakhs  11.183367
0                age   4.567634
1  number_of_dependants  4.534650
4      insurance_plan   3.584752
10  marital_status_Unmarried  3.411185
8      region_Southeast   2.922414
5      total_risk_score   2.687610
9      region_Southwest   2.670666
6      gender_Male       2.421496
16  employment_status_Salaried  2.382134
17  employment_status_Self-Employed  2.137753
7      region_Northwest   2.102556
15  smoking_status_Regular   1.777089
12  bmi_category_Overweight   1.549922
11  bmi_category_Obesity     1.352806
13  bmi_category_Underweight   1.302886
14  smoking_status_Occasional   1.272745

```

After the Variance Inflation Factor (VIF) was calculated for all features, it was observed that both `income_level` and `income_lakhs` had VIF scores exceeding the commonly accepted threshold of 10, indicating the presence of high multicollinearity between them.



To address this issue and enhance model stability, the `income_level` feature was dropped, as it exhibited the higher VIF value of the two. This step was taken to reduce redundancy without significantly compromising the information contained in the dataset.

```
[88]: # Initialize a list to store highly collinear features to be removed later
```

```
high_vif_features = []
```

```
[89]: # Identify the feature with the highest VIF score and add it to the removal_
      ↪ list 'high_vif_features'
```

```
highest_vif_feature = vif_df['features'][vif_df['vif_score'].idxmax()]
high_vif_features.append(highest_vif_feature)
print(high_vif_features)
```

```
['income_level']
```

```
[90]: # Drop the most collinear feature from the temporary DataFrame
```

```
temp_df1 = temp_df.drop(['income_level'],axis=1).copy()
temp_df1
```

```
[90]:
```

	age	number_of_dependants	income_lakhs	insurance_plan	\
0	0.148148	0.0	0.050505	0.0	
1	0.203704	0.4	0.050505	0.0	
2	0.574074	0.4	0.191919	0.5	
3	0.222222	0.6	0.767677	1.0	
4	0.000000	0.0	0.989899	0.5	
...	...	...	...	...	
49903	0.111111	0.0	0.343434	0.0	
49904	0.537037	0.4	0.818182	1.0	
49905	0.055556	0.0	0.313131	0.0	
49906	0.000000	0.4	0.191919	0.0	
49907	0.555556	0.6	0.060606	0.5	

	total_risk_score	gender_Male	region_Northwest	region_Southeast	\
0	0.428571	1	1	0	
1	0.428571	0	0	1	
2	0.428571	0	0	0	
3	0.000000	0	0	1	
4	0.428571	1	0	0	
...	...	...	...	...	
49903	0.000000	0	1	0	
49904	0.357143	0	0	1	
49905	0.000000	1	1	0	
49906	0.000000	1	1	0	
49907	0.428571	0	0	0	

	region_Southwest	marital_status_Unmarried	bmi_category_Obesity	\
0	0	1	0	
1	0	0	1	
2	0	0	0	
3	0	0	0	
4	0	1	0	
...	...	...	...	
49903	0	1	0	
49904	0	0	0	
49905	0	1	0	
49906	0	1	0	
49907	1	0	0	

	bmi_category_Overweight	bmi_category_Underweight	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	1	0	
...	...	...	
49903	0	1	
49904	0	0	
49905	0	0	
49906	0	0	
49907	0	0	

	smoking_status_Occasional	smoking_status_Regular	\
0	0	0	
1	0	1	
2	0	0	
3	0	0	
4	0	1	
...	...	...	
49903	0	0	
49904	0	0	
49905	0	1	
49906	0	0	
49907	1	0	

	employment_status_Salaried	employment_status_Self-Employed
0	1	0
1	1	0
2	0	1
3	1	0
4	0	1
...	...	...
49903	0	1

49904	1	0
49905	0	0
49906	1	0
49907	0	1

[49908 rows x 17 columns]

The same process will be repeated until all VIF values are within the acceptable threshold.

```
[91]: # Recalculate VIF scores for the updated feature set

vif_dict = {'features': [], 'vif_score': []}

for i,col in enumerate(temp_df1.columns):
    vif = variance_inflation_factor(temp_df1,i)
    vif_dict['features'].append(col)
    vif_dict['vif_score'].append(vif)

# Create a DataFrame from the updated VIF scores
vif_df = pd.DataFrame(vif_dict)

# Sort VIF scores in descending order to identify next candidates for removal
vif_df.sort_values(by='vif_score',ascending=False)
```

```
[91]:
```

	features	vif_score
0	age	4.545825
1	number_of_dependants	4.526598
3	insurance_plan	3.445682
9	marital_status_Unmarried	3.393718
7	region_Southeast	2.919775
4	total_risk_score	2.687326
8	region_Southwest	2.668314
2	income_lakhs	2.480563
5	gender_Male	2.409980
15	employment_status_Salaried	2.374628
16	employment_status_Self-Employed	2.132810
6	region_Northwest	2.100789
14	smoking_status_Regular	1.777024
11	bmi_category_Overweight	1.549907
10	bmi_category_Obesity	1.352748
12	bmi_category_Underweight	1.302636
13	smoking_status_Occasional	1.272744

All remaining features have acceptable VIF scores.

```
[92]: # Remove the previously identified high-VIF features from the original
      ↪ DataFrame.
final_df = df7.drop(high_vif_features,axis=1)
```

```
# Display the final cleaned dataset
final_df
```

```
[92]:
```

	age	number_of_dependants	income_lakhs	insurance_plan	\
0	0.148148	0.0	0.050505	0.0	
1	0.203704	0.4	0.050505	0.0	
2	0.574074	0.4	0.191919	0.5	
3	0.222222	0.6	0.767677	1.0	
4	0.000000	0.0	0.989899	0.5	
...	...	...	...	...	
49903	0.111111	0.0	0.343434	0.0	
49904	0.537037	0.4	0.818182	1.0	
49905	0.055556	0.0	0.313131	0.0	
49906	0.000000	0.4	0.191919	0.0	
49907	0.555556	0.6	0.060606	0.5	

	annual_premium_amount	total_risk_score	gender_Male	region_Northwest	\
0	9053	0.428571	1	1	
1	16339	0.428571	0	0	
2	18164	0.428571	0	0	
3	20303	0.000000	0	0	
4	13365	0.428571	1	0	
...	...	...	...	...	
49903	9111	0.000000	0	1	
49904	27076	0.357143	0	0	
49905	8564	0.000000	1	1	
49906	9490	0.000000	1	1	
49907	19730	0.428571	0	0	

	region_Southeast	region_Southwest	marital_status_Unmarried	\
0	0	0	1	
1	1	0	0	
2	0	0	0	
3	1	0	0	
4	0	0	1	
...	...	...	...	
49903	0	0	1	
49904	1	0	0	
49905	0	0	1	
49906	0	0	1	
49907	0	1	0	

	bmi_category_Obesity	bmi_category_Overweight	\
0	0	0	
1	1	0	
2	0	0	

3	0	0
4	0	1
...	...	...
49903	0	0
49904	0	0
49905	0	0
49906	0	0
49907	0	0

	bmi_category_Underweight	smoking_status_Occasional \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...	...	...
49903	1	0
49904	0	0
49905	0	0
49906	0	0
49907	0	1

	smoking_status_Regular	employment_status_Salaried \
0	0	1
1	1	1
2	0	0
3	0	1
4	1	0
...	...	...
49903	0	0
49904	0	1
49905	1	0
49906	0	1
49907	0	0

	employment_status_Self-Employed
0	0
1	0
2	1
3	0
4	1
...	...
49903	1
49904	0
49905	0
49906	0
49907	1

[49908 rows x 18 columns]

## 0.5 Model Selection

The model selection process proceeded as follows:

1. Split the dataset into training and test sets.
2. Trained different models on the training data.
3. Used cross-validation to compare model performance.
4. Tuned hyperparameters of the best model.
5. Evaluated the final model's performance on the test set.

### 0.5.1 Dataset Split

The dataset is divided as follows:

- Training set: 70%
- Test set: 30%

Split was done randomly with a fixed seed for reproducibility.

```
[93]: # Separate features and target variable

features = final_df.drop(['annual_premium_amount'],axis=1)
target = final_df['annual_premium_amount']

[94]: # Split the data into training and test sets

X_train,X_test,y_train,y_test = train_test_split(features,target,test_size=0.
↪3,random_state=42)

[95]: # Display the shape of training and test sets

print(f"X_train shape: {X_train.shape}, y_train shape: {y_train.shape}")
print(f"X_test shape: {X_test.shape}, y_test shape: {y_test.shape}")
```

```
X_train shape: (34935, 17), y_train shape: (34935,)
X_test shape: (14973, 17), y_test shape: (14973,)
```

### 0.5.2 Model Training

Various models will be tried and training will be performed, including:

- Linear Regression

- Ridge Regression
- Lasso Regression
- Random Forest Regressor
- XGBoost Regressor

## Linear Regression

```
[96]: # Initialize and train the Linear Regression model
lr=LinearRegression()
lr.fit(X_train,y_train)

# Evaluate model performance
train_score = lr.score(X_train,y_train)
test_score = lr.score(X_test,y_test)

# Print the R2 scores
print(f'Train Score : {train_score} , Test Score : {test_score}')
```

Train Score : 0.9280957176093705 , Test Score : 0.9283765993531427

```
[97]: # Predict on test data
y_pred = lr.predict(X_test)

# Calculate Mean Squared Error and Root Mean Squared Error
mse = mean_squared_error(y_test,y_pred)
rmse = root_mean_squared_error(y_test,y_pred)

# Print performance metrics
print(f'MSE : {mse} , RMSE : {rmse}')
```

MSE : 5056639.130347778 , RMSE : 2248.697207350909

```
[98]: # Feature Names and Its Coefficients

print('Features -> ',lr.feature_names_in_)
print('\n','**'*50)
print('Coef -> ',lr.coef_)
```

```
Features -> ['age' 'number_of_dependants' 'income_lakhs' 'insurance_plan'
'total_risk_score' 'gender_Male' 'region_Northwest' 'region_Southeast'
'region_Southwest' 'marital_status_Unmarried' 'bmi_category_Obesity'
'bmi_category_Overweight' 'bmi_category_Underweight'
'smoking_status_Occasional' 'smoking_status_Regular'
'employment_status_Salaried' 'employment_status_Self-Employed']
```

```
*****
*****
```

```
Coef -> [11236.4082633  -536.61172956  -353.97067084 12515.43782972
 4846.88626348   121.53023031  -34.50219135   27.97504587
 -23.47372614  -821.78088646  3356.07750448  1613.1421694
 364.31350886   722.41051531  2261.97816483   149.57179545
 378.25130285]
```

```
[99]: # Creating a dataframe of features and coefficients
```

```
feat_coef = {
    'features' : lr.feature_names_in_,
    'coef' : lr.coef_
}

feat_coef_df = pd.DataFrame(feat_coef)
feat_coef_df.sort_values(by=['coef'],ascending=False)
```


```
[99]:
```

	features	coef
3	insurance_plan	12515.437830
0	age	11236.408263
4	total_risk_score	4846.886263
10	bmi_category_Obesity	3356.077504
14	smoking_status_Regular	2261.978165
11	bmi_category_Overweight	1613.142169
13	smoking_status_Occasional	722.410515
16	employment_status_Self-Employed	378.251303
12	bmi_category_Underweight	364.313509
15	employment_status_Salaried	149.571795
5	gender_Male	121.530230
7	region_Southeast	27.975046
8	region_Southwest	-23.473726
6	region_Northwest	-34.502191
2	income_lakhs	-353.970671
1	number_of_dependants	-536.611730
9	marital_status_Unmarried	-821.780886

```
[100]: # Plotting the features and its coefficients
```

```
plt.figure(figsize=(12,10))
sns.barplot(data=feat_coef_df.sort_values(by=['coef'],ascending=False),x =_
    ↪ 'coef',y = 'features' )
plt.title('Coefficient of Each Feature')
plt.xlabel('Coefficient')
plt.ylabel('Features')
plt.show()
```





## Lasso Regression

```
[101]: # Initialize and train the Lasso Regression model
ls = Lasso()
ls.fit(X_train,y_train)

# Evaluate model performance
train_score = ls.score(X_train,y_train)
test_score = ls.score(X_test,y_test)

# Print the  $R^2$  scores
print(f'Train Score : {train_score} , Test Score : {test_score}')
```

Train Score : 0.9280927085624446 , Test Score : 0.9283637752728616

```
[102]: # Predict on test data
y_pred = ls.predict(X_test)

# Calculate Mean Squared Error and Root Mean Squared Error
mse = mean_squared_error(y_test,y_pred)
rmse = root_mean_squared_error(y_test,y_pred)

# Print performance metrics
```

```
print(f'MSE : {mse} , RMSE : {rmse}')
```

MSE : 5057544.515257937 , RMSE : 2248.8985115513633

### Ridge Regression

```
[103]: # Initialize and train the Ridge Regression model
rg = Ridge()
rg.fit(X_train,y_train)

# Evaluate model performance
train_score = rg.score(X_train,y_train)
test_score = rg.score(X_test,y_test)

# Print the  $R^2$  scores
print(f'Train Score : {train_score} , Test Score : {test_score}')
```

Train Score : 0.9280956798900618 , Test Score : 0.9283764905001947

```
[104]: # Predict on test data
y_pred = rg.predict(X_test)

# Calculate Mean Squared Error and Root Mean Squared Error
mse = mean_squared_error(y_test,y_pred)
rmse = root_mean_squared_error(y_test,y_pred)

# Print performance metrics
print(f'MSE : {mse} , RMSE : {rmse}')
```

MSE : 5056646.815407011 , RMSE : 2248.698916130617

### Observation

Neither Lasso nor Ridge regression showed significant improvement over Linear Regression.

### Random Forest Regressor

```
[105]: # Initialize and train the Random Forest Regression model
rfr = RandomForestRegressor()
rfr.fit(X_train,y_train)

# Evaluate model performance
train_score = rfr.score(X_train,y_train)
test_score = rfr.score(X_test,y_test)

# Print the  $R^2$  scores
print(f'Train Score : {train_score} , Test Score : {test_score}')
```

Train Score : 0.9965657768248808 , Test Score : 0.9789939664693872

```
[106]: # Predict on test data
y_pred = rfr.predict(X_test)

# Calculate Mean Squared Error and Root Mean Squared Error
mse = mean_squared_error(y_test,y_pred)
rmse = root_mean_squared_error(y_test,y_pred)

# Print performance metrics
print(f'MSE : {mse} , RMSE : {rmse}')
```

MSE : 1483033.8990467205 , RMSE : 1217.798792513246

Based on the  $R^2$  score, it can be seen that this model fits the data well. The XGBoost model will be tried next.

### XGboost Regressor

```
[107]: # Initialize and train the XGboost model
xgb = XGBRegressor()
xgb.fit(X_train,y_train)

# Evaluate model performance
train_score = xgb.score(X_train,y_train)
test_score = xgb.score(X_test,y_test)

# Print the  $R^2$  scores
print(f'Train Score : {train_score} , Test Score : {test_score} ')
```

Train Score : 0.9861041903495789 , Test Score : 0.9807721972465515

```
[108]: # Predict on test data
y_pred = xgb.predict(X_test)

# Calculate Mean Squared Error and Root Mean Squared Error
mse = mean_squared_error(y_test,y_pred)
rmse = root_mean_squared_error(y_test,y_pred)

# Print performance metrics
print(f'MSE : {mse} , RMSE : {rmse}')
```

MSE : 1357488.75 , RMSE : 1165.1131591796875

### Model Performance Comparison: XGBoost vs. Random Forest

Both XGBoost and Random Forest models are observed to perform well on the dataset. However, the following points are noted:

- A lower training score and a higher test score are yielded by the XGBoost model compared to Random Forest.
- This suggests that the XGBoost model generalizes better and is less prone to overfitting.
- To validate this observation, cross-validation will be used to evaluate both models more robustly using `cross_validate`.

**Cross Validation - RF vs XGboost** Stratified K-Fold will not be performed since this is a regression task; therefore, K-Fold cross-validation is considered sufficient.

```
[109]: # Set up 5-fold cross-validation

kf = KFold(n_splits = 5,shuffle=True,random_state=42)
```

Cross-validation will be performed on both Random Forest and XGBoost models, with the run time also being recorded.

The `cross_validate` function is used instead of `cross_val_score` to obtain both training and testing scores.

```
[110]: # Perform cross-validation for Random Forest Regressor

start_time = time()
cv_rfr =
    ↪cross_validate(rfr,features,target,cv=kf,return_train_score=True,scoring='r2')
end_time = time()
total_time_rfr = end_time - start_time
print(f'Total Time Taken : {round(total_time_rfr,2)} seconds')
```

Total Time Taken : 67.49 seconds

```
[111]: # Display training scores
print("Training scores (R²):", cv_rfr['train_score'])

# Display test scores
print("Validation scores (R²):", cv_rfr['test_score'])
```

Training scores (R²): [0.9965103 0.99652913 0.99644755 0.99654474 0.99645785]  
Validation scores (R²): [0.97901928 0.97902788 0.97945362 0.97865826 0.9794153 ]

```
[112]: # Perform cross-validation for XGboost Regressor

start_time = time()
cv_xgb =
    ↪cross_validate(xgb,features,target,cv=kf,return_train_score=True,scoring='r2')
end_time = time()
total_time_xgb = end_time - start_time
print(f'Total Time Taken : {round(total_time_xgb,2)} seconds')
```

Total Time Taken : 0.74 seconds

```
[113]: # Display training scores
print("Training scores (R²):", cv_xgb['train_score'])

# Display test scores
print("Validation scores (R²):", cv_xgb['test_score'])
```

Training scores ( $R^2$ ): [0.98580199 0.9857831 0.98563534 0.98585421 0.98580426]  
Validation scores ( $R^2$ ): [0.9809615 0.98077649 0.98098087 0.98112887 0.98120141]

```
[114]: cv_df = pd.DataFrame(
    {
        'model' : ['Random Forest','XGboost'],
        'execution_time (seconds)' : [total_time_rfr,total_time_xgb],
        'average_train_score' : [np.mean(cv_rfr['train_score']),np.
↪mean(cv_xgb['train_score'])],
        'average_test_score' : [np.mean(cv_rfr['test_score']),np.
↪mean(cv_xgb['test_score'])]
    }
)

cv_df
```

```
[114]:
```

	model	execution_time (seconds)	average_train_score \
0	Random Forest	67.486237	0.996498
1	XGboost	0.744500	0.985776

	average_test_score
0	0.979115
1	0.981010

### Cross-Validation Summary: XGBoost vs. Random Forest

The cross-validation results confirm that XGBoost is consistently observed to generalize better than the Random Forest model:

- *Lower training scores and higher test scores* were achieved by XGBoost, indicating *better generalization* and *reduced overfitting*.
- In terms of performance, XGBoost was approximately **83 times faster** than Random Forest during cross-validation.

Given its superior predictive performance and computational efficiency, **XGBoost will be selected as the final model.**